

O'REILLY®



# Learning Spark

LIGHTNING-FAST DATA ANALYSIS

Holden Karau,  
Andy Kowinski & Matei Zaharia

# Learning Spark

**Holden Karau**

**Andy Konwinski**

**Patrick Wendell**

**Matei Zaharia**



Beijing • Cambridge • Farnham • Köln • Sebastopol • Tokyo

## Table of Contents

<b>Preface.....</b>	<b>5</b>
Audience.....	5
How This Book is Organized.....	6
Supporting Books.....	6
Code Examples.....	7
Early Release Status and Feedback.....	7
<b>Chapter 1. Introduction to Data Analysis with Spark.....</b>	<b>8</b>
What is Apache Spark?.....	8
A Unified Stack.....	8
Who Uses Spark, and For What?.....	11
A Brief History of Spark.....	13
Spark Versions and Releases.....	13
Spark and Hadoop.....	14
<b>Chapter 2. Downloading and Getting Started.....</b>	<b>15</b>
Downloading Spark.....	15
Introduction to Spark's Python and Scala Shells.....	16
Introduction to Core Spark Concepts.....	20
Standalone Applications.....	23
Conclusion.....	25
<b>Chapter 3. Programming with RDDs.....</b>	<b>26</b>
RDD Basics.....	26
Creating RDDs.....	28
RDD Operations.....	28
Passing Functions to Spark.....	32
Common Transformations and Actions.....	36
Persistence (Caching).....	46
Conclusion.....	48
<b>Chapter 4. Working with Key-Value Pairs.....</b>	<b>49</b>

Motivation.....	49
Creating Pair RDDs.....	49
Transformations on Pair RDDs.....	50
Actions Available on Pair RDDs.....	60
Data Partitioning.....	61
Conclusion.....	70
<b>Chapter 5. Loading and Saving Your Data.....</b>	<b>71</b>
Motivation.....	71
Choosing a Format.....	71
Formats.....	72
File Systems.....	88
Compression.....	89
Databases.....	91
Conclusion.....	93
<b>About the Authors.....</b>	<b>95</b>

## Preface

As parallel data analysis has become increasingly common, practitioners in many fields have sought easier tools for this task. Apache Spark has quickly emerged as one of the most popular tools for this purpose, extending and generalizing MapReduce. Spark offers three main benefits. First, it is easy to use—you can develop applications on your laptop, using a high-level API that lets you focus on the content of your computation. Second, Spark is fast, enabling interactive use and complex algorithms. And third, Spark is a *general* engine, allowing you to combine multiple types of computations (e.g., SQL queries, text processing and machine learning) that might previously have required learning different engines. These features make Spark an excellent starting point to learn about big data in general.

This introductory book is meant to get you up and running with Spark quickly. You'll learn how to learn how to download and run Spark on your laptop and use it interactively to learn the API. Once there, we'll cover the details of available operations and distributed execution. Finally, you'll get a tour of the higher-level libraries built into Spark, including libraries for machine learning, stream processing, graph analytics and SQL. We hope that this book gives you the tools to quickly tackle data analysis problems, whether you do so on one machine or hundreds.

## Audience

This book targets Data Scientists and Engineers. We chose these two groups because they have the most to gain from using Spark to expand the scope of problems they can solve. Spark's rich collection of data focused libraries (like MLlib) make it easy for data scientists to go beyond problems that fit on single machine while making use of their statistical background. Engineers, meanwhile, will learn how to write general-purpose distributed programs in Spark and operate production applications. Engineers and data scientists will both learn different details from this book, but will both be able to apply Spark to solve large distributed problems in their respective fields.

Data scientists focus on answering questions or building models from data. They often have a statistical or math background and some familiarity with tools like Python, R and SQL. We have made sure to include Python, and wherever possible SQL, examples for all our material, as well as an overview of the machine learning and advanced analytics libraries in Spark. If you are a data scientist, we hope that after reading this book you will be able to use the same mathematical approaches to solving problems, except much faster and on a much larger scale.

The second group this book targets is software engineers who have some experience with Java, Python or another programming language. If you are an engineer, we hope that this book will show you how to set up a Spark cluster, use the Spark shell, and write Spark applications to solve parallel processing problems. If you are familiar with Hadoop, you have a bit of a head start on figuring out how to interact with HDFS and how to manage a cluster, but either way, we will cover basic distributed execution concepts.

Regardless of whether you are a data analyst or engineer, to get the most of this book you should have some familiarity with one of Python, Java, Scala, or a similar language. We assume that you already have a solution for storing your data and we cover how to load and save data from many common ones, but not how to set them up. If you don't have experience with one of those languages, don't worry, there are excellent resources available to learn these. We call out some of the books available in [Supporting Books](#).

## **How This Book is Organized**

The chapters of this book are laid out in such a way that you should be able to go through the material front to back. At the start of each chapter, we will mention which sections of the chapter we think are most relevant to data scientists and which sections we think are most relevant for engineers. That said, we hope that all the material is accessible to readers of either background.

The first two chapters will get you started with getting a basic Spark installation on your laptop and give you an idea of what you can accomplish with Apache Spark. Once we've got the motivation and setup out of the way, we will dive into the Spark Shell, a very useful tool for development and prototyping. Subsequent chapters then cover the Spark programming interface in detail, how applications execute on a cluster, and higher-level libraries available on Spark such as Spark SQL and MLlib.

## **Supporting Books**

If you are a data scientist and don't have much experience with Python, the Learning Python book is an excellent introduction.

If you are an engineer and after reading this book you would like to expand your data analysis skills, Machine Learning for Hackers and Doing Data Science are excellent books from O'Reilly.

This book is intended to be accessible to beginners. We do intend to release a deep dive follow-up for those looking to gain a more thorough understanding of Spark's internals.

## Code Examples

All of the code examples found in this book are on GitHub. You can examine them and check them out from <https://github.com/databricks/learning-spark>. Code examples are provided in Java, Scala, and Python.

### Tip

Our Java examples are written to work with Java version 6 and higher. Java 8 introduces a new syntax called “lambdas” that makes writing inline functions much easier, which can simplify Spark code. We have chosen not to take advantage of this syntax in most of our examples, as most organizations are not yet using Java 8. If you would like to try Java 8 syntax, you can see [the Databricks blog post on this topic](#).

## Early Release Status and Feedback

This is an **early release** copy of Learning Spark, and as such we are still working on the text, adding code examples, and writing some of the later chapters. Although we hope that the book is useful in its current form, we would greatly appreciate your feedback so we can improve it and make the best possible finished product. The authors and editors can be reached at [book-feedback@databricks.com](mailto:book-feedback@databricks.com).

The authors would like to thank the reviewers who offered feedback so far: Juliet Hougland, Andrew Gal, Michael Gregson, Stephan Jou, Josh Mahonin, and Mike Patterson.

# Chapter 1. Introduction to Data Analysis with Spark

This chapter provides a high level overview of what Apache Spark is. If you are already familiar with Apache Spark and its components, feel free to jump ahead to [Chapter 2](#).

## What is Apache Spark?

Apache Spark is a cluster computing platform designed to be *fast* and *general-purpose*.

On the speed side, Spark extends the popular MapReduce model to efficiently support more types of computations, including interactive queries and stream processing. Speed is important in processing large datasets as it means the difference between exploring data interactively and waiting minutes between queries, or waiting hours to run your program versus minutes. One of the main features Spark offers for speed is the ability to run computations in memory, but the system is also faster than MapReduce for complex applications running on disk.

On the generality side, Spark is designed to cover a wide range of workloads that previously required separate distributed systems, including batch applications, iterative algorithms, interactive queries and streaming. By supporting these workloads in the same engine, Spark makes it easy and inexpensive to *combine* different processing types, which is often necessary in production data analysis pipelines. In addition, it reduces the management burden of maintaining separate tools.

Spark is designed to be highly accessible, offering simple APIs in Python, Java, Scala and SQL, and rich built-in libraries. It also integrates closely with other big data tools. In particular, Spark can run in Hadoop clusters and access any Hadoop data source.

## A Unified Stack

The Spark project contains multiple closely-integrated components. At its core, Spark is a “computational engine” that is responsible for scheduling, distributing, and monitoring applications consisting of many computational tasks across many worker machines, or a *computing cluster*. Because the core engine of Spark is both fast and general-purpose, it powers multiple higher-level components specialized for various workloads, such as SQL or machine learning. These components are designed to interoperate closely, letting you combine them like libraries in a software project.



A philosophy of tight integration has several benefits. First, all libraries and higher level components in the stack benefit from improvements at the lower layers. For example, when Spark's core engine adds an optimization, SQL and machine learning libraries automatically speed up as well. Second, the costs associated with running the stack are minimized, because instead of running 5-10 independent software systems, an organization only needs to run one. These costs include deployment, maintenance, testing, support, and more. This also means that each time a new component is added to the Spark stack, every organization that uses Spark will immediately be able to try this new component. This changes the cost of trying out a new type of data analysis from downloading, deploying, and learning a new software project to upgrading Spark.

Finally, one of the largest advantages of tight integration is the ability to build applications that seamlessly combine different processing models. For example, in Spark you can write one application that uses machine learning to classify data in real time as it is ingested from streaming sources. Simultaneously analysts can query the resulting data, also in real-time, via SQL, e.g. to join the data with unstructured log files. In addition, more sophisticated data engineers can access the same data via the Python shell for ad-hoc analysis. Others might access the data in standalone batch applications. All the while, the IT team only has to maintain one software stack.

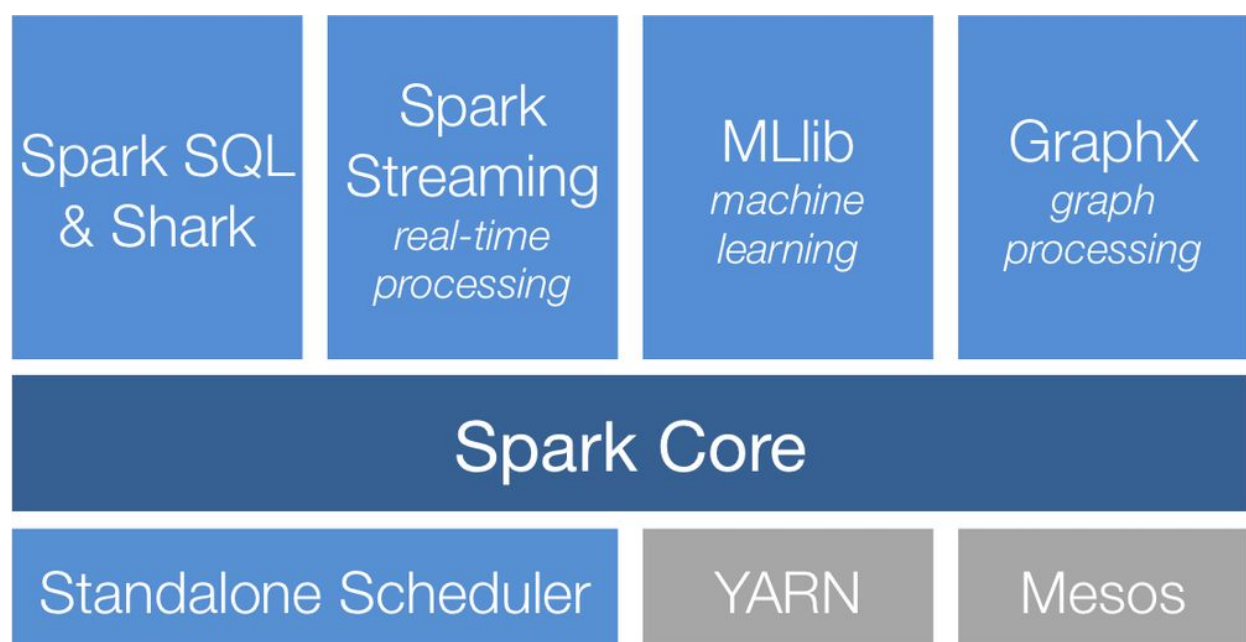


Figure 1-1. The Spark Stack

Here we will briefly introduce each of the components shown in [Figure 1-1](#).

## **Spark Core**

Spark Core contains the basic functionality of Spark, including components for task scheduling, memory management, fault recovery, interacting with storage systems, and more. Spark Core is also home to the API that defines Resilient Distributed Datasets (RDDs), which are Spark's main programming abstraction. RDDs represent a collection of items distributed across many compute nodes that can be manipulated in parallel. Spark Core provides many APIs for building and manipulating these collections.

## **Spark SQL**

Spark SQL provides support for interacting with Spark via SQL as well as the Apache Hive variant of SQL, called the Hive Query Language (HiveQL). Spark SQL represents database tables as Spark RDDs and translates SQL queries into Spark operations. Beyond providing the SQL interface to Spark, Spark SQL allows developers to intermix SQL queries with the programmatic data manipulations supported by RDDs in Python, Java and Scala, all within a single application. This tight integration with the rich and sophisticated computing environment provided by the rest of the Spark stack makes Spark SQL unlike any other open source data warehouse tool. Spark SQL was added to Spark in version 1.0.

Shark is a project out of UC Berkeley that pre-dates Spark SQL and is being ported to work on top of Spark SQL. Shark provides additional functionality so that Spark can act as drop-in replacement for Apache Hive. This includes a HiveQL shell, as well as a JDBC server that makes it easy to connect external graphing and data exploration tools.

## **Spark Streaming**

Spark Streaming is a Spark component that enables processing live streams of data. Examples of data streams include log files generated by production web servers, or queues of messages containing status updates posted by users of a web service. Spark Streaming provides an API for manipulating data streams that closely matches the Spark Core's RDD API, making it easy for programmers to learn the project and move between applications that manipulate data stored in memory, on disk, or arriving in real-time. Underneath its API, Spark Streaming was designed to provide the same degree of fault tolerance, throughput, and scalability that the Spark Core provides.

## **MLlib**

Spark comes with a library containing common machine learning (ML) functionality called MLlib. MLlib provides multiple types of machine learning algorithms, including binary classification, regression, clustering and collaborative filtering, as well as supporting

functionality such as model evaluation and data import. It also provides some lower level ML primitives including a generic gradient descent optimization algorithm. All of these methods are designed to scale out across a cluster.

## GraphX

GraphX is a library added in Spark 0.9 that provides an API for manipulating graphs (e.g., a social network's friend graph) and performing graph-parallel computations. Like Spark Streaming and Spark SQL, GraphX extends the Spark RDD API, allowing us to create a directed graph with arbitrary properties attached to each vertex and edge. GraphX also provides set of operators for manipulating graphs (e.g., subgraph and mapVertices) and a library of common graph algorithms (e.g., PageRank and triangle counting).

## Cluster Managers

Under the hood, Spark is designed to efficiently scale up from one to many thousands of compute nodes. To achieve this while maximizing flexibility, Spark can run over a variety of *cluster managers*, including including Hadoop YARN, Apache Mesos, and a simple cluster manager included in Spark itself called the Standalone Scheduler. If you are just installing Spark on an empty set of machines, the Standalone Scheduler provides an easy way to get started; while if you already have a Hadoop YARN or Mesos cluster, Spark's support for these allows your applications to also run on them.

## Who Uses Spark, and For What?

Because Spark is a general purpose framework for cluster computing, it is used for a diverse range of applications. In the [Preface](#) we outlined two personas that this book targets as readers: Data Scientists and Engineers. Let's take a closer look at each of these personas and how they use Spark. Unsurprisingly, the typical use cases differ across the two personas, but we can roughly classify them into two categories, *data science* and *data applications*.

Of course, these are imprecise personas and usage patterns, and many folks have skills from both, sometimes playing the role of the investigating Data Scientist, and then "changing hats" and writing a hardened data processing system. Nonetheless, it can be illuminating to consider the two personas and their respective use cases separately.

## Data Science Tasks

Data Science is the name of a discipline that has been emerging over the past few years centered around analyzing data. While there is no standard definition, for our purposes a *Data Scientist*

is somebody whose main task is to analyze and model data. Data scientists may have experience using SQL, statistics, predictive modeling (machine learning), and some programming, usually in Python, Matlab or R. Data scientists also have experience with techniques necessary to transform data into formats that can be analyzed for insights (sometimes referred to as *data wrangling*).

Data Scientists use their skills to analyze data with the goal of answering a question or discovering insights. Oftentimes, their workflow involves ad-hoc analysis, and so they use interactive shells (vs. building complex applications) that let them see results of queries and snippets of code in the least amount of time. Spark's speed and simple APIs shine for this purpose, and its built-in libraries mean that many algorithms are available out of the box.

Sometimes, after the initial exploration phase, the work of a Data Scientist will be "productionized", or extended, hardened (i.e. made fault tolerant), and tuned to become a production data processing application, which itself is a component of a business application. For example, the initial investigation of a Data Scientist might lead to the creation of a production recommender system that is integrated into on a web application and used to generate customized product suggestions to users. Often it is a different person or team that leads the process of productizing the work of the Data Scientists, and that person is often an Engineer.

## **Data Processing Applications**

The other main use case of Spark can be described in the context of the *Engineer* persona. For our purposes here, we think of Engineers as large class of software developers who use Spark to build production data processing applications. These developers usually have an understanding of the principles of software engineering, such as encapsulation, interface design, and Object Oriented Programming. They frequently have a degree in Computer Science. They use their engineering skills to design and build software systems that implement a business use case.

For Engineers, Spark provides a simple way to parallelize these applications across clusters, and hides the complexity of distributed systems programming, network communication and fault tolerance. The system gives enough control to monitor, inspect and tune applications while allowing common tasks to be implemented quickly. The modular nature of the API (based on passing distributed collections of objects) makes it easy to factor work into reusable libraries and test it locally.

Spark's users choose to use it for their data processing applications because it provides a wide variety of functionality, is easy to learn and use, and is mature and reliable.

## A Brief History of Spark

Spark is an open source project that has been built and is maintained by a thriving and diverse community of developers from many different organizations. If you or your organization are trying Spark for the first time, you might be interested in the history of the project. Spark started in 2009 as a research project in the UC Berkeley RAD Lab, later to become the AMPLab. The researchers in the lab had previously been working on Hadoop MapReduce, and observed that MapReduce was inefficient for iterative and interactive computing jobs. Thus, from the beginning, Spark was designed to be fast for interactive queries and iterative algorithms, bringing in ideas like support for in-memory storage and efficient fault recovery.

Research papers were published about Spark at academic conferences and soon after its creation in 2009, it was already 10–20x faster than MapReduce for certain jobs.

Some of Spark's first users were other groups inside of UC Berkeley, including machine learning researchers such as the the Mobile Millennium project, which used Spark to monitor and predict traffic congestion in the San Francisco bay Area. In a very short time, however, many external organizations began using Spark, and today, over 50 organizations list themselves on the [Spark PoweredBy page](#) <sup>[1]</sup>, and dozens speak about their use cases at Spark community events such as [Spark Meetups](#) <sup>[2]</sup> and the [Spark Summit](#) <sup>[3]</sup>. Apart from UC Berkeley, major contributors to the project currently include Yahoo!, Intel and Databricks.

In 2011, the AMPLab started to develop higher-level components on Spark, such as Shark (Hive on Spark) and Spark Streaming. These and other components are often referred to as the [Berkeley Data Analytics Stack \(BDAS\)](#) <sup>[4]</sup>. BDAS includes both components of Spark and other software projects that complement it, such as the Tachyon memory manager.

Spark was first open sourced in March 2010, and was transferred to the Apache Software Foundation in June 2013, where it is now a top-level project.

## Spark Versions and Releases

Since its creation Spark has been a very active project and community, with the number of contributors growing with each release. Spark 1.0 had over 100 individual contributors. Though the level of activity has rapidly grown, the community continues to release updated versions of Spark on a regular schedule. Spark 1.0 was released in May 2014. This book focuses primarily on Spark 1.0 and beyond, though most of the concepts and examples also work in earlier versions.

# Spark and Hadoop

Spark can create distributed datasets from any file stored in the Hadoop distributed file system (HDFS) or other storage systems supported by Hadoop (including your local file system, Amazon S3, Cassandra, Hive, HBase, etc). Spark supports text files, SequenceFiles, Avro, Parquet, and any other Hadoop InputFormat. We will look at interacting with these data sources in the [loading and saving chapter](#).

---

[1] <https://cwiki.apache.org/confluence/display/SPARK/Powered+By+Spark>

[2] <http://www.meetup.com/spark-users/>

[3] <http://spark-summit.org>

[4] <https://amplab.cs.berkeley.edu/software>

## Chapter 2. Downloading and Getting Started

In this chapter we will walk through the process of downloading and running Spark in local mode on a single computer. This chapter was written for anybody that is new to Spark, including both Data Scientists and Engineers.

Spark can be used from Python, Java or Scala. To benefit from this book, you don't need to be an expert programmer, but we do assume that you are comfortable with the basic syntax of at least one of these languages. We will include examples in all languages wherever possible.

Spark itself is written in Scala, and runs on the Java Virtual Machine (JVM). To run Spark on either your laptop or a cluster, all you need is an installation of Java 6 (or newer). If you wish to use the Python API you will also need a Python interpreter (version 2.6 or newer) . Spark does not yet work with Python 3.

### Downloading Spark

The first step to using Spark is to download and unpack it into a usable form. Let's start by downloading a recent precompiled released version of Spark. Visit <http://spark.apache.org/downloads.html>, then under “Pre-built packages”, next to “For Hadoop 1 (HDP1, CDH3)”, click “direct file download”. This will download a compressed tar file, or “tarball,” called `spark-1.0.0-bin-hadoop1.tgz`.

If you want to use Spark with another Hadoop version, those are also available from <http://spark.apache.org/downloads.html> but will have slightly different file names. Building from source is also possible, and you can find the latest source code on GitHub at <http://github.com/apache/spark>.

### Note

Most Unix and Linux variants, including Mac OS X, come with a command-line tool called `tar` that can be used to unpack tar files. If your operating system does not have the `tar` command installed, try searching the Internet for a free tar extractor—for example, on Windows, you may wish to try 7-Zip.

Now that we have downloaded Spark, let's unpack it and take a look at what comes with the default Spark distribution. To do that, open a terminal, change to the directory where you downloaded Spark, and `untar` the file. This will create a new directory with the same name but

without the final .tgz suffix. Change into the directory, and see what's inside. You can use the following commands to accomplish all of that.

```
cd ~  
tar -xf spark-1.0.0-bin-hadoop1.tgz  
cd spark-1.0.0-bin-hadoop1  
ls
```

In the line containing the tar command above, the x flag tells tar we are extracting files, and the f flag specifies the name of the tarball. The ls command lists the contents of the Spark directory. Let's briefly consider the names and purpose of some of the more important files and directories you see here that come with Spark.

- **README.md** - Contains short instructions for getting started with Spark.
- **bin** - Contains executable files that can be used to interact with Spark in various ways, e.g. the spark-shell, which we will cover later in this chapter, is in here.
- **core, streaming, python** - source code of major components of the Spark project.
- **examples** - contains some helpful Spark standalone jobs that you can look at and run to learn about the Spark API.

Don't worry about the large number of directories and files the Spark project comes with; we will cover most of these in the rest of this book. For now, let's dive in right away and try out Spark's Python and Scala shells. We will start by running some of the examples that come with Spark. Then we will write, compile and run a simple Spark Job of our own.

All of the work we will do in this chapter will be with Spark running in "local mode", i.e. non-distributed mode, which only uses a single machine. Spark can run in a variety of different modes, or environments. Beyond local mode, Spark can also be run on Mesos, YARN, on top of a Standalone Scheduler that is included in the Spark distribution. We will cover the various deployment modes in detail in chapter (to come).

## Introduction to Spark's Python and Scala Shells

Spark comes with interactive shells that make ad-hoc data analysis easy. Spark's shells will feel familiar if you have used other shells such as those in R, Python, and Scala, or operating system shells like Bash or the Windows command prompt.

Unlike most other shells, however, which let you manipulate data using the disk and memory on a single machine, Spark's shells allow you to interact with data that is distributed on disk or in



memory across many machines, and Spark takes care of automatically distributing this processing.

Because Spark can load data into memory, many distributed computations, even ones that process terabytes of data across dozens of machines, can finish running in a few seconds. This makes the sort of iterative, ad-hoc, and exploratory analysis commonly done in shells a good fit for Spark. Spark provides both Python and Scala shells that have been augmented to support connecting to a cluster.

## Note

Most of this book includes code in all of Spark's languages, but interactive shells are only available in Python and Scala. Because a shell is very useful for learning the API, we recommend using one of these languages for these examples even if you are a Java developer. The API is the same in every language.

The easiest way to demonstrate the power of Spark's shells is to start using one of them for some simple data analysis. Let's walk through the example from the Quick Start Guide in the official Spark documentation [\[5\]](#).

The first step is to open up one of Spark's shells. To open the Python version of the Spark Shell, which we also refer to as the PySpark Shell, go into your Spark directory and type:

```
bin/pyspark
```

(Or `bin\pyspark` in Windows.) To open the Scala version of the shell, type:

```
bin/spark-shell
```

The shell prompt should appear within a few seconds. When the shell starts, you will notice a lot of log messages. You may need to hit [Enter] once to clear the log output, and get to a shell prompt. Figure [Figure 2-1](#) shows what the PySpark shell looks like when you open it.





**Distributed Datasets**, or **RDDs**. RDDs are Spark's fundamental abstraction for distributed data and computation.

Before we say more about RDDs, let's create one in the shell from a local text file and do some very simple ad-hoc analysis by following the example below.

Example 2-1. Python line count

```
>>> lines = sc.textFile("README.md") # Create an RDD called lines

>>> lines.count() # Count the number of items in this RDD
127

>>> lines.first() # First item in this RDD, i.e. first line of README.md
u'# Apache Spark'
```

Example 2-2. Scala line count

```
scala> val lines = sc.textFile("README.md") // Create an RDD called lines
lines: spark.RDD[String] = MappedRDD[...]

scala> lines.count() // Count the number of items in this RDD
res0: Long = 127

scala> lines.first() // First item in this RDD, i.e. first line of README.md
res1: String = # Apache Spark
```

To exit the shell, you can press Control+D.

In the example above, the variables called `lines` are RDDs, created here from a text file on our local machine. We can run various parallel operations on the RDDs, such as counting the number of elements in the dataset (here lines of text in the file) or printing the first one. We will discuss RDDs in great depth in later chapters, but before we go any further, let's take a moment now to introduce basic Spark concepts.

## Introduction to Core Spark Concepts

Now that you have run your first Spark code using the shell, it's time learn about programming in it in more detail.

At a high level, every Spark application consists of a *driver program* that launches various parallel operations on a cluster. The driver program contains your application's main function and defines distributed datasets on the cluster, then applies operations to them. In the examples

above, the driver program was the Spark shell itself, and you could just type in the operations you wanted to run.

Driver programs access Spark through a `SparkContext` object, which represents a connection to a computing cluster. In the shell, a `SparkContext` is automatically created for you, as the variable called `sc`. Try printing out `sc` to see its type:

```
>>> sc
<pyspark.context.SparkContext object at 0x1025b8f90>
```

Once you have a `SparkContext`, you can use it to build *resilient distributed datasets*, or *RDDs*. In the example above, we called `SparkContext.textFile` to create an RDD representing the lines of text in a file. We can then run various operations on these lines, such as `count()`.

To run these operations, driver programs typically manage a number of nodes called *executors*. For example, if we were running the `count()` above on a cluster, different machines might count lines in different ranges of the file. Because we just ran the Spark shell locally, it executed all its work on a single machine—but you can connect the same shell to a cluster to analyze data in parallel. [Figure 2-3](#) shows how Spark executes on a cluster.




Figure 2-3. Components for distributed execution in Spark

Finally, a lot of Spark’s API revolves around passing functions to its operators to run them on the cluster. For example, we could extend our README example by *filtering* the lines in the file that contain a word, such as “Python”:

Example 2-3. Python filtering example

```
>>> lines = sc.textFile("README.md")

>>> pythonLines = lines.filter(lambda line: "Python" in line)

>>> pythonLines.first()
```

```
u'## Interactive Python Shell'
```

Example 2-4. Scala filtering example

```
scala> val lines = sc.textFile("README.md") // Create an RDD called lines
lines: spark.RDD[String] = MappedRDD[...]
```

```
scala> val pythonLines = lines.filter(line => line.contains("Python"))
pythonLines: spark.RDD[String] = FilteredRDD[...]
```

```
scala> lines.first()
reso: String = ## Interactive Python Shell
```

## Note

If you are unfamiliar with the lambda or => syntax above, it is a shorthand way to define functions inline in Python and Scala. When using Spark in these languages, you can also define a function separately and then pass its name to Spark. For example, in Python:

```
def hasPython(line):
    return "Python" in line
```

```
pythonLines = lines.filter(hasPython)
```

Passing functions to Spark is also possible in Java, but in this case they are defined as classes, implementing an interface called Function. For example:

```
JavaRDD<String> pythonLines = lines.filter(
    new Function<String, Boolean>() {
        Boolean call(String line) { return line.contains("Python"); }
    }
);
```

Java 8 introduces shorthand syntax called “lambdas” that looks similar to Python and Scala. Here is how the code would look with this syntax:

```
JavaRDD<String> pythonLines = lines.filter(line -> line.contains("Python"));
```

We discuss passing functions further in [Passing Functions to Spark](#).

While we will cover the Spark API in more detail later, a lot of its magic is that function-based operations like `filter` *also* parallelize across the cluster. That is, Spark automatically takes your function (e.g. `line.contains("Python")`) and ships it to executor nodes. Thus, you can write code in a single driver program and automatically have parts of it run on multiple nodes. [Chapter 3](#) covers the RDD API in more detail.

## Standalone Applications

The final piece missing in this quick tour of Spark is how to use it in standalone programs. Apart from running interactively, Spark can be linked into standalone applications in either Java, Scala or Python. The main difference from using it in the shell is that you need to initialize your own `SparkContext`. After that, the API is the same.

The process of linking to Spark varies by language. In Java and Scala, you give your application a Maven dependency on the `spark-core` artifact published by Apache. As of the time of writing, the latest Spark version is 1.0.0, and the Maven coordinates for that are:

```
groupId = org.apache.spark  
artifactId = spark-core_2.10  
version = 1.0.0
```

If you are unfamiliar with Maven, it is a popular package management tool for Java-based languages that lets you link to libraries in public repositories. You can use Maven itself to build your project, or use other tools that can talk to the Maven repositories, including Scala's SBT tool or Gradle. Popular integrated development environments like Eclipse also allow you to directly add a Maven dependency to a project.

In Python, you simply write applications as Python scripts, but you must instead run them using a special `bin/spark-submit` script included in Spark. This script sets up the environment for Spark's Python API to function. Simply run your script with:

```
bin/spark-submit my_script.py
```

(Note that you will have to use backslashes instead of forward slashes on Windows.)

## Note

In Spark versions before 1.0, use `bin/pyspark my_script.py` to run Python applications instead.

For detailed examples of linking applications to Spark, refer to the [Quick Start Guide](#) <sup>[6]</sup> in the official Spark documentation. In a final version of the book, we will also include full examples in an appendix.

## Initializing a SparkContext

Once you have linked an application to Spark, you need to import the Spark packages in your program and create a `SparkContext`. This is done by first creating a `SparkConf` object to configure your application, and then building a `SparkContext` for it. Here is a short example in each supported language:

Example 2-5. Initializing Spark in Python

```
from pyspark import SparkConf, SparkContext

conf = SparkConf().setMaster("local").setAppName("My App")
sc = SparkContext(conf)
```

Example 2-6. Initializing Spark in Java

```
import org.apache.spark.SparkConf;
import org.apache.spark.api.java.JavaSparkContext;

SparkConf conf = new SparkConf().setMaster("local").setAppName("My App");
JavaSparkContext sc = new JavaSparkContext(conf);
```

Example 2-7. Initializing Spark in Scala

```
import org.apache.spark.SparkConf
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

val conf = new SparkConf().setMaster("local").setAppName("My App")
val sc = new SparkContext("local", "My App")
```

These examples show the minimal way to initialize a `SparkContext`, where you pass two parameters:

- A *cluster URL*, namely “local” in these examples, which tells Spark how to connect to a cluster. “local” is a special value that runs Spark on one thread on the local machine, without connecting to a cluster.



- An *application name*, namely “My App” in these examples. This will identify your application on the cluster manager’s UI if you connect to a cluster.

Additional parameters exist for configuring how your application executes or adding code to be shipped to the cluster, but we will cover these in later chapters of the book.

After you have initialized a `SparkContext`, you can use all the methods we showed before to create RDDs (e.g. from a text file) and manipulate them.

Finally, to shut down Spark, you can either call the `stop()` method on your `SparkContext`, or simply exit the application (e.g. with `System.exit(0)` or `sys.exit()`).

This quick overview should be enough to let you run a standalone Spark application on your laptop. For more advanced configuration, a later chapter in the book will cover how to connect your application to a cluster, including packaging your application so that its code is automatically shipped to worker nodes. For now, please refer to the [Quick Start Guide](http://spark.apache.org/docs/latest/quick-start.html) <sup>[7]</sup> in the official Spark documentation.

## Conclusion

In this chapter, we have covered downloading Spark, running it locally on your laptop, and using it either interactively or from a standalone application. We gave a quick overview of the core concepts involved in programming with Spark: a driver program creates a `SparkContext` and RDDs, and then runs parallel operations on them. In the next chapter, we will dive more deeply into how RDDs operate.

---

<sup>[5]</sup> <http://spark.apache.org/docs/latest/quick-start.html>

<sup>[6]</sup> <http://spark.apache.org/docs/latest/quick-start.html>

<sup>[7]</sup> <http://spark.apache.org/docs/latest/quick-start.html>

## Chapter 3. Programming with RDDs

This chapter introduces Spark’s core abstraction for working with data, the Resilient Distributed Dataset (RDD). An RDD is simply a distributed collection of elements. In Spark all work is expressed as either creating new RDDs, transforming existing RDDs, or calling operations on RDDs to compute a result. Under the hood, Spark automatically distributes the data contained in RDDs across your cluster and parallelizes the operations you perform on them.

Both Data Scientists and Engineers should read this chapter, as RDDs are the core concept in Spark. We highly recommend that you try some of these examples in an interactive shell (see [Introduction to Spark’s Python and Scala Shells](#)). In addition, all code in this chapter is available in the book’s [GitHub repository](#).

### RDD Basics

An RDD in Spark is simply a distributed collection of objects. Each RDD is split into multiple *partitions*, which may be computed on different nodes of the cluster. RDDs can contain any type of Python, Java or Scala objects, including user-defined classes.

Users create RDDs in two ways: by loading an external dataset, or by distributing a collection of objects in their driver program. We have already seen loading a text file as an RDD of strings using `SparkContext.textFile()`:

Example 3-1. Creating an RDD of strings with `textFile()` in Python

```
>>> lines = sc.textFile("README.md")
```

Once created, RDDs offer two types of operations: *transformations* and *actions*. *Transformations* construct a new RDD from a previous one. For example, one transformation we saw before is filtering data that matches a predicate. In our text file example, we can use this to create a new RDD holding just the strings that contain “Python”:

```
>>> pythonLines = lines.filter(lambda line: "Python" in line)
```

*Actions*, on the other hand, compute a result based on an RDD, and either return it to the driver program or save it to an external storage system (e.g., HDFS). One example of an action we called earlier is `first()`, which returns the first element in an RDD:

```
>>> pythonLines.first()
```

```
u'## Interactive Python Shell'
```

The difference between transformations and actions is due to the way Spark computes RDDs. Although you can define new RDDs any time, Spark only computes them in a *lazy* fashion, the first time they are used in an action. This approach might seem unusual at first, but makes a lot of sense when working with big data. For instance, consider the example above, where we defined a text file and then filtered the lines with “Python”. If Spark were to load and store all the lines in the file as soon as we wrote `lines = sc.textFile(...)`, it would waste a lot of storage space, given that we then immediately filter out many lines. Instead, once Spark sees the whole chain of transformations, it can compute just the data needed for its result. In fact, for the `first()` action, Spark only scans the file until it finds the first matching line; it doesn’t even read the whole file.

Finally, Spark’s RDDs are by default recomputed each time you run an action on them. If you would like to reuse an RDD in multiple actions, you can ask Spark to *persist* it using `RDD.persist()`. After computing it the first time, Spark will store the RDD contents in memory (partitioned across the machines in your cluster), and reuse them in future actions. Persisting RDDs on disk instead of memory is also possible. The behavior of not persisting by default may again seem unusual, but it makes a lot of sense for big datasets: if you will not reuse the RDD, there’s no reason to waste storage space when Spark could instead stream through the data once and just compute the result.<sup>[8]</sup>

In practice, you will often use `persist` to load a subset of your data into memory and query it repeatedly. For example, if we knew that we wanted to compute multiple results about the README lines that contain “Python”, we could write:

```
>>> pythonLines.persist()
```

```
>>> pythonLines.count()
```

```
2
```

```
>>> pythonLines.first()
```

```
u'## Interactive Python Shell'
```

To summarize, every Spark program and shell session will work as follows:

1. Create some input RDDs from external data.
2. Transform them to define new RDDs using transformations like `filter()`.
3. Ask Spark to `persist()` any intermediate RDDs that will need to be reused.

4. Launch actions such as `count()` and `first()` to kick off a parallel computation, which is then optimized and executed by Spark.

In the rest of this chapter, we'll go through each of these steps in detail, and cover some of the most common RDD operations in Spark.

## Creating RDDs

Spark provides two ways to create RDDs: loading an external dataset and parallelizing a collection in your driver program.

The simplest way to create RDDs is to take an existing in-memory collection and pass it to `SparkContext`'s `parallelize` method. This approach is very useful when learning Spark, since you can quickly create your own RDDs in the shell and perform operations on them. Keep in mind however, that outside of prototyping and testing, this is not widely used since it requires you have your entire dataset in memory on one machine.

Example 3-2. Python `parallelize` example

```
lines = sc.parallelize(["pandas", "i like pandas"])
```

Example 3-3. Scala `parallelize` example

```
val lines = sc.parallelize(List("pandas", "i like pandas"))
```

Example 3-4. Java `parallelize` example

```
JavaRDD<String> lines = sc.parallelize(Arrays.asList("pandas", "i like pandas"));
```

A more common way to create RDDs is to load data in external storage. Loading external datasets is covered in detail in [Chapter 5](#). However, we already saw one method that loads a text file as an RDD of strings, `SparkContext.textFile`:

Example 3-5. Python `textFile` example

```
lines = sc.textFile("/path/to/README.md")
```

Example 3-6. Scala `textFile` example

```
val lines = sc.textFile("/path/to/README.md")
```

Example 3-7. Java `textFile` example

```
JavaRDD<String> lines = sc.textFile("/path/to/README.md");
```

## RDD Operations

RDDs support two types of operations, *transformations* and *actions*. Transformations are operations on RDDs that return a new RDD, such as `map` and `filter`. Actions are operations that return a result back to the driver program or write it to storage, and kick off a computation, such

as count and first. Spark treats transformations and actions very differently, so understanding which type of operation you are performing will be important. If you are ever confused whether a given function is a transformation or an action, you can look at its return type: transformations return RDDs whereas actions return some other data type.

## Transformations

Transformations are operations on RDDs that return a new RDD. As discussed in the lazy evaluation section, transformed RDDs are computed lazily, only when you use them in an action. Many transformations are element-wise, that is they work on one element at a time, but this is not true for all transformations.

As an example, suppose that we have a log file, log.txt, with a number of messages, and we want to select only the error messages. We can use the filter transformation seen before. This time though, we'll show a filter in all three of Spark's language APIs:

Example 3-8. Python filter example

```
inputRDD = sc.textFile("log.txt")
errorsRDD = inputRDD.filter(lambda x: "error" in x)
```

Example 3-9. Scala filter example

```
val inputRDD = sc.textFile("log.txt")
val errorsRDD = inputRDD.filter(line => line.contains("error"))
```

Example 3-10. Java filter example

```
JavaRDD<String> inputRDD = sc.textFile("log.txt");
JavaRDD<String> errorsRDD = inputRDD.filter(
    new Function<String, Boolean>() {
        public Boolean call(String x) { return x.contains("error");
    }
});
```

Note that the filter operation does not mutate the existing inputRDD. Instead, it returns a pointer to an entirely new RDD. inputRDD can still be re-used later in the program, for instance, to search for other words. In fact, let's use inputRDD again to search for lines with the word "warning" in them. Then, we'll use another transformation, union, to print out the number of lines that contained either "error" or "warning". We show Python here, but the union() function is identical in all three languages:

Example 3-11. Python union example

```
errorsRDD = inputRDD.filter(lambda x: "error" in x)
warningsRDD = inputRDD.filter(lambda x: "warning" in x)
```

```
badLinesRDD = errorsRDD.union(warningsRDD)
```

union is a bit different than filter, in that it operates on two RDDs instead of one. Transformations can actually operate on any number of input RDDs.

Finally, as you derive new RDDs from each other using transformations, Spark keeps track of the set of dependencies between different RDDs, called the *lineage graph*. It uses this information to compute each RDD on demand and to recover lost data if part of a persistent RDD is lost. We will show a lineage graph for this example in [Figure 3-1](#).

## Actions

We've seen how to create RDDs from each other with transformations, but at some point, we'll want to actually *do something* with our dataset. Actions are the second type of RDD operation. They are the operations that return a final value to the driver program or write data to an external storage system. Actions force the evaluation of the transformations required for the RDD they are called on, since they are required to actually produce output.

Continuing the log example from the previous section, we might want to print out some information about the badLinesRDD. To do that, we'll use two actions, count(), which returns the count as a number, and take(), which collects a number of elements from the RDD.

Example 3-12. Python error count example using actions

```
print "Input had " + badLinesRDD.count() + " concerning lines"
print "Here are 10 examples:"
for line in badLinesRDD.take(10):
    print line
```

Example 3-13. Scala error count example using actions

```
println("Input had " + badLinesRDD.count() + " concerning lines")
println("Here are 10 examples:")
badLinesRDD.take(10).foreach(println)
```

Example 3-14. Java error count example using actions

```
System.out.println("Input had " + badLinesRDD.count() + " concerning lines")
System.out.println("Here are 10 examples:")
for (String line: badLinesRDD.take(10)) {
    System.out.println(line);
}
```

In this example, we used take() to retrieve a small number of elements in the RDD at the driver program. We then iterate over them locally to print out information at the driver. RDDs also

have a `collect()` function to retrieve the entire RDD. This can be useful if your program filters RDDs down to a very small size and you'd like to deal with it locally. Keep in mind that your entire dataset must fit in memory on a single machine to use `collect()` on it, so `collect()` shouldn't be used on large datasets.

In most cases RDDs can't just be `collect()`'ed to the driver because they are too large. In these cases, it's common to write data out to a distributed storage systems such as HDFS or Amazon S3. The contents of an RDD can be saved using the `saveAsTextFile` action, `saveAsSequenceFile` or any of a number of actions for various built-in formats. We will cover the different options for exporting data later on in [Chapter 5](#).

The image below presents the lineage graph for this entire example, starting with our `inputRDD` and ending with the two actions. It is important to note that each time we call a new action, the entire RDD must be computed "from scratch". To avoid this inefficiency, users can *persist* intermediate results, as we will cover in [Persistence \(Caching\)](#).

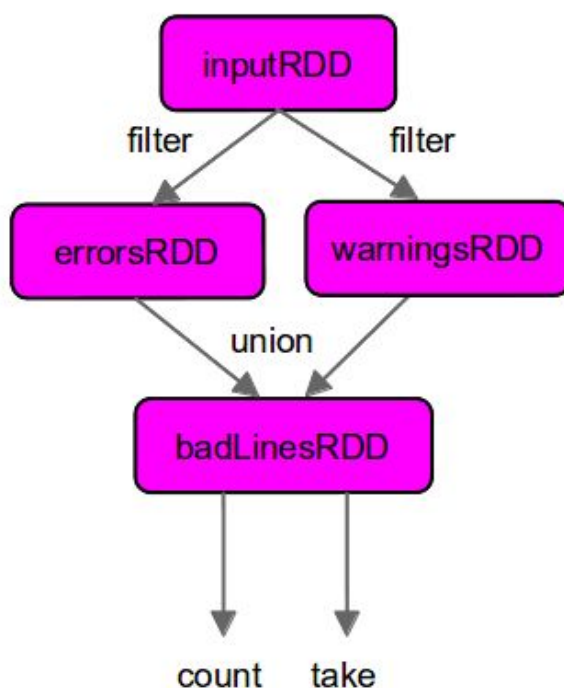


Figure 3-1. RDD lineage graph created during log analysis.

## Lazy Evaluation

Transformations on RDDs are lazily evaluated, meaning that Spark will not begin to execute until it sees an action. This can be somewhat counter-intuitive for new users, but may be familiar for those who have used functional languages such as Haskell or LINQ-like data processing frameworks.

Lazy evaluation means that when we call a transformation on an RDD (for instance calling `map`), the operation is not immediately performed. Instead, Spark internally records meta-data to indicate this operation has been requested. Rather than thinking of an RDD as containing specific data, it is best to think of each RDD as consisting of instructions on how to compute the data that we build up through transformations. Loading data into an RDD is lazily evaluated in the same way transformations are. So when we call `sc.textFile` the data is not loaded until it is necessary. Like with transformations, the operation (in this case reading the data) can occur multiple times.

## Tip

Although transformations are lazy, force Spark to execute them at any time by running an action, such as `count()`. This is an easy way to test out just part of your program.

Spark uses lazy evaluation to reduce the number of passes it has to take over our data by grouping operations together. In MapReduce systems like Hadoop, developers often have to spend a lot of time considering how to group together operations to minimize the number of MapReduce passes. In Spark, there is no substantial benefit to writing a single complex map instead of chaining together many simple operations. Thus, users are free to organize their program into smaller, more manageable operations.

## Passing Functions to Spark

Most of Spark transformations, and some of its actions, depend on passing in functions that are used by Spark to compute data. Each of the core languages has a slightly different mechanism for passing functions to Spark.

## Python

In Python, we have three options for passing functions into Spark. For shorter function we can pass in lambda expressions, as we have done in the example at the start of this chapter. We can also pass in top-level functions, or locally defined functions.

Example 3-15. Passing a lambda in Python

```
word = rdd.filter(lambda s: "error" in s)
```

### Passing a top-level Python function.

```
def containsError(s):  
    return "error" in s
```



```
word = rdd.filter(containsError)
```

One issue to watch out for when passing functions is that if you pass functions that are members of an object, or references to fields in an object (e.g., `self.field`), this results in sending in the *entire* object, which can be much larger than just the bit of information you need. Sometimes this can also cause your program to fail, if your class contains objects that Python can't figure out how to pickle.

Example 3-16. Passing a function with field references (don't do this!)

```
class SearchFunctions(object):
    def __init__(self, query):
        self.query = query
    def isMatch(self, s):
        return query in s
    def getMatchesFunctionReference(self, rdd):
        # Problem: references all of "self" in "self.isMatch"
        return rdd.filter(self.isMatch)
    def getMatchesMemberReference(self, rdd):
        # Problem: references all of "self" in "self.query"
        return rdd.filter(lambda x: self.query in x)
```

Instead, just extract the fields you need from your object into local variable and pass that in, like we do below:

Example 3-17. Python function passing without field references

```
class WordFunctions(object):
    ...
    def getMatchesNoReference(self, rdd):
        # Safe: extract only the field we need into a local variable
        query = self.query
        return rdd.filter(lambda x: query in x)
```

## Scala

In Scala, we can pass in functions defined inline or references to methods or static functions as we do for Scala's other functional APIs. Some other considerations come into play though, namely that the function we pass and the data referenced in it needs to be `Serializable` (implementing Java's `Serializable` interface). Furthermore, like in Python, passing a method or field of an object includes a reference to that whole object, though this is less obvious because we are not forced to write these references with `self`. Like how we did with Python, we can

instead extract out the fields we need as local variables and avoid needing to pass the whole object containing them.

Example 3-18. Scala function passing

```
class SearchFunctions(val query: String) {  
  def isMatch(s: String): Boolean = {  
    s.contains(query)  
  }  
  def getMatchesFunctionReference(rdd: RDD[String]): RDD[String] = {  
    // Problem: "isMatch" means "this.isMatch", so we pass all of "this"  
    rdd.map(isMatch)  
  }  
  def getMatchesFieldReference(rdd: RDD[String]): RDD[String] = {  
    // Problem: "query" means "this.query", so we pass all of "this"  
    rdd.map(x => x.split(query))  
  }  
  def getMatchesNoReference(rdd: RDD[String]): RDD[String] = {  
    // Safe: extract just the field we need into a local variable  
    val query_ = this.query  
    rdd.map(x => x.split(query_))  
  }  
}
```

If you “NotSerializableException” errors in Scala, a reference to a method or field in a non-serializable class is usually the problem. Note that passing in local variables or functions that are members of a top-level object is always safe.

## Java

In Java, functions are specified as objects that implement [one of Spark’s function interfaces](#) from the `org.apache.spark.api.java.function` package. There are a number of different interfaces based on the return type of the function. We show the most basic function interfaces below, and cover a number of [other function interfaces](#) for when we need to return special types of data in the [section on converting between RDD types](#).

Table 3-1. Standard Java function interfaces

Function name	method to implement	Usage
Function<T, R>	R call(T)	Take in one input and return one output, for use with things like map and filter.
Function2<T1, T2, R>	R call(T1, T2)	Take in two inputs and return one output, for use with things like aggregate or fold.
FlatMapFunction<T, R>	Iterable<R> call(T)	Take in one input and return zero or more outputs, for use with things like flatMap.

We can either define our function classes in-line as anonymous inner classes, or make a named class:

Example 3-19. Java function passing with anonymous inner class

```
RDD<String> errors = lines.filter(new Function<String, Boolean>() {
    public Boolean call(String x) { return s.contains("error"); }
});
```

Example 3-20. Java function passing with named class

```
class ContainsError implements Function<String, Boolean>() {
    public Boolean call(String x) { return x.contains("error"); }
}
```

```
RDD<String> errors = lines.filter(new ContainsError());
```

The style to choose is a personal preference, but we find that top-level named functions are often cleaner for organizing large programs. One other benefit of top-level functions is that you can give them constructor parameters:

Example 3-21. Java function class with parameters

```
class Contains implements Function<String, Boolean>() {
    private String query;
    public Contains(String query) { this.query = query; }
    public Boolean call(String x) { return x.contains(query); }
}
```

```
RDD<String> errors = lines.filter(new Contains("error"));
```

In Java 8, you can also use lambda expressions to concisely implement the Function interfaces. Since Java 8 is still relatively new as of the writing of this book, our examples use the more verbose syntax for defining classes in previous versions of Java. However, with lambda expressions, our search example would look like this:

Example 3-22. Java function passing with lambda expression in Java 8

```
RDD<String> errors = lines.filter(s -> s.contains("error"));
```

If you are interested in using Java 8's lambda expression, refer to [Oracle's documentation](#) and [the Databricks blog post on how to use lambdas with Spark](#).

## Tip

Both anonymous inner classes and lambda expressions can reference any final variables in the method enclosing them, so you can pass these variables to Spark just like in Python and Scala.

## Common Transformations and Actions

In this chapter, we tour the most common transformations and actions in Spark. Additional operations are available on RDDs containing certain type of data—for example, statistical functions on RDDs of numbers, and key-value operations such as aggregating data by key on RDDs of key-value pairs. We cover converting between RDD types and these special operations in later sections.

### Basic RDDs

We will begin by evaluating what operations we can do on all RDDs regardless of the data. These transformations and actions are available on all RDD classes.

## Transformations

### Element-wise transformations

The two most common transformations you will likely be performing on basic RDDs are map, and filter. The map transformation takes in a function and applies it to each element in the RDD with the result of the function being the new value of each element in the resulting RDD. The filter transformation take in a function and returns an RDD which only has elements that pass the filter function.




Figure 3-2. Map and filter on an RDD

We can use map to do any number of things from fetching the website associated with each URL in our collection, to just squaring the numbers. With Scala and python you can use the standard anonymous function notation or pass in a function, and with Java you should use Spark's Function class from `org.apache.spark.api.java.function` or Java 8 functions.

It is useful to note that the return type of the map does not have to be the same as the input type, so if we had an RDD of customer IDs and our map function were to fetch the corresponding customer records the type of our input RDD would be `RDD[CustomerID]` and the type of the resulting RDD would be `RDD[CustomerRecord]`.

Lets look at a basic example of map which squares all of the numbers in an RDD:

Example 3-23. Python squaring the value in an RDD

```

nums = sc.parallelize([1, 2, 3, 4])
squared = nums.map(lambda x: x * x).collect()
for num in squared:
    print "%i " % (num)
  
```

Example 3-24. Scala squaring the values in an RDD

```

val input = sc.parallelize(List(1, 2, 3, 4))
val result = input.map(x => x*x)
println(result.collect())
  
```

Example 3-25. Java squaring the values in an RDD

```

JavaRDD<Integer> rdd = sc.parallelize(Arrays.asList(1, 2, 3, 4));
JavaRDD<Integer> result = rdd.map(new Function<Integer, Integer>() {
    public Integer call(Integer x) { return x*x; }
});
System.out.println(StringUtils.join(result.collect(), ","));
  
```

Sometimes we want to produce multiple output elements for each input element. The operation to do this is called `flatMap`. Like with `map`, the function we provide to `flatMap` is called

individually for each element in our input RDD. Instead of returning a single element, we return an iterator with our return values. Rather than producing an RDD of iterators, we get back an RDD which consists of the elements from all of the iterators. A simple example of flatMap is splitting up an input string into words, as shown below.

Example 3-26. Python flatMap example, splitting lines into words

```
lines = sc.parallelize(["hello world", "hi"])
words = lines.flatMap(lambda line: line.split(" "))
words.first() # returns "hello"
```

Example 3-27. Scala flatMap example, splitting lines into multiple words

```
val lines = sc.parallelize(List("hello world", "hi"))
val words = lines.flatMap(line => line.split(" "))
words.first() // returns "hello"
```

Example 3-28. Scala flatMap example, splitting lines into multiple words

```
JavaRDD<String> lines = sc.parallelize(Arrays.asList("hello world", "hi"));
JavaRDD<String> words = rdd.flatMap(new FlatMapFunction<String, String>() {
    public Iterable<String> call(String line) {
        return Arrays.asList(line.split(" "));
    }
});
words.first(); // returns "hello"
```

## Pseudo Set Operations

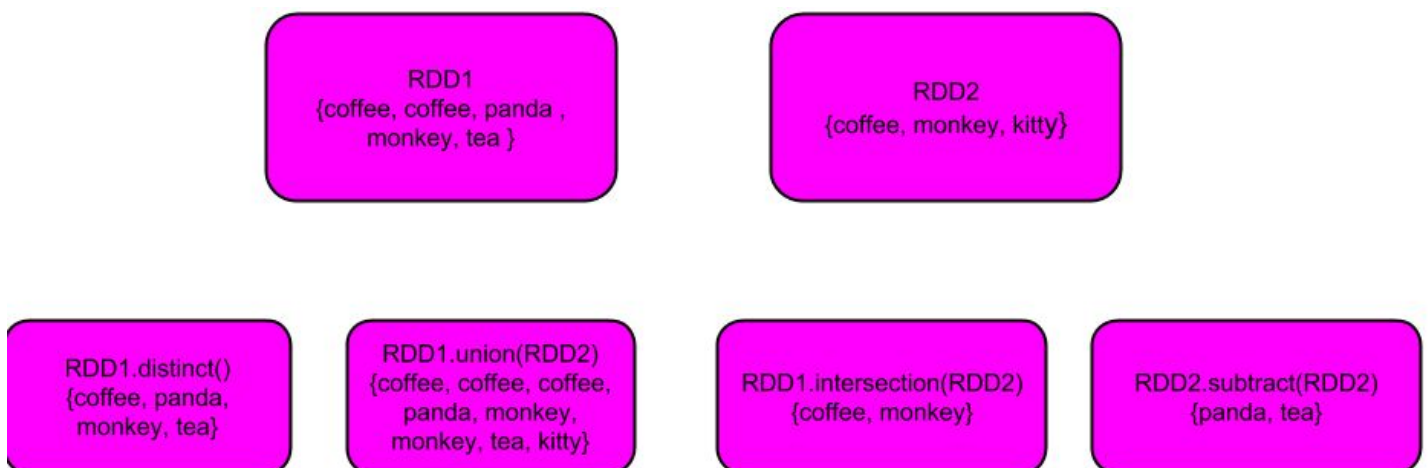


Figure 3-3. Some simple set operations (image to be redone)

RDDs support many of the operations of mathematical sets, such as union and intersection, even when the RDDs themselves not properly sets.

The set property most frequently missing from our RDDs is the uniqueness of elements. If we only want unique elements we can use the `RDD.distinct()` transformation to produce a new RDD with only distinct items. Note that `distinct()` is expensive, however, as it requires shuffling all the data over the network to ensure that we only receive one copy of each element.

The simplest set operation is `union(other)`, which gives back an RDD consisting of the data from both sources. This can be useful in a number of use cases, such as processing log files from many sources. Unlike the mathematical `union()`, if there are duplicates in the input RDDs, the result of Spark's `union()` will contain duplicates (which we can fix if desired with `distinct()`).

Spark also provides an `intersection(other)` method, which returns only elements in both RDDs. `intersection()` also removes all duplicates (including duplicates from a single RDD) while running. While intersection and union are to very similar concepts, the performance of intersection is much worse since it requires a shuffle over the network to identify common elements.

Sometimes we need to remove some data from consideration. The `subtract(other)` function takes in another RDD and returns an RDD that only has values present in the first RDD and not the second RDD.

We can also compute a Cartesian product between two RDDs. The `cartesian(other)` transformation results in possible pairs of (a, b) where a is in the source RDD and b is in the other RDD. The Cartesian product can be useful when we wish to consider the similarity between all possible pairs, such as computing every users expected interests in each offer. We can also take the Cartesian product of an RDD with itself, which can be useful for tasks like computing user similarity.

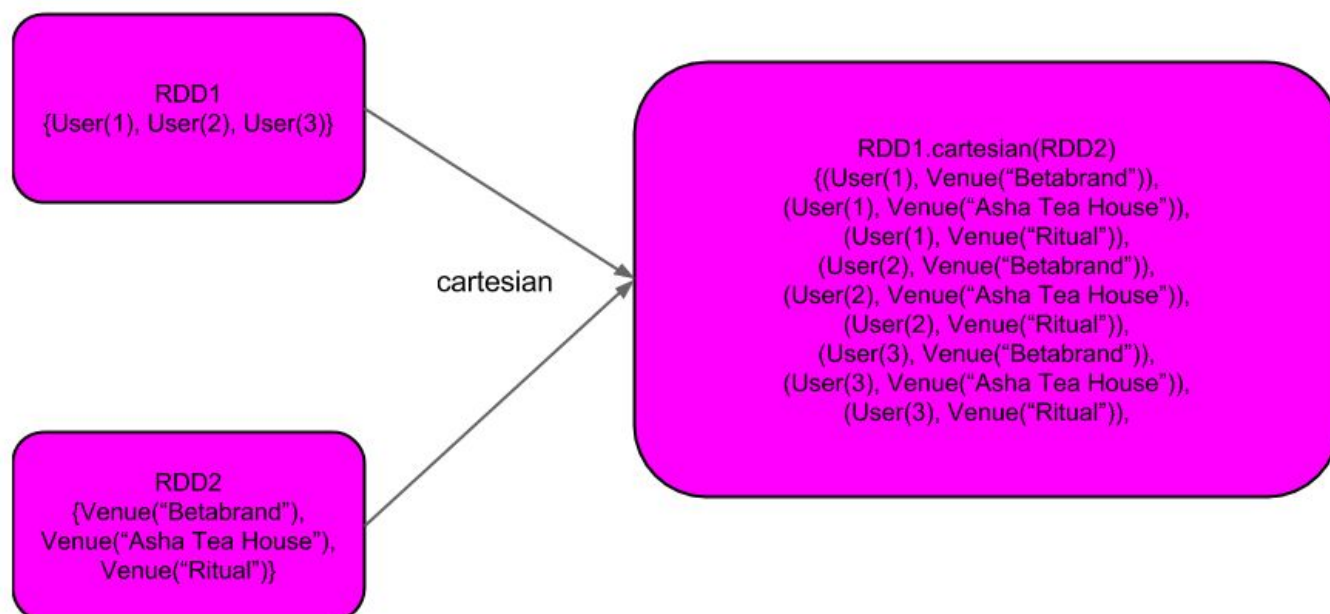


Figure 3-4. Cartesian product between two RDDs

The tables below summarize common single-RDD and multi-RDD transformations.

Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}

Function Name	Purpose	Example	Result
map	Apply a function to each element in the RDD and return an RDD of the result	<code>rdd.map(x =&gt; x + 1)</code>	{2, 3, 4, 4}
flatMap	Apply a function to each element in the RDD and return an RDD of the contents of the iterators returned. Often used to extract words.	<code>rdd.flatMap(x =&gt; x.to(3))</code>	{1, 2, 3, 2, 3, 3, 3}
filter	Return an RDD consisting of only elements which pass the condition passed to filter	<code>rdd.filter(x =&gt; x != 1)</code>	{2, 3, 3}
distinct	Remove duplicates	<code>rdd.distinct()</code>	{1, 2, 3}
sample(withReplacement, fraction, [seed])	Sample an RDD	<code>rdd.sample(false, 0.5)</code>	non-deterministic

Table 3-3. Two-RDD transformations on RDDs containing {1, 2, 3} and {3, 4, 5}

Function Name	Purpose	Example	Result
union	Produce an RDD contain elements from both RDDs	<code>rdd.union(other)</code>	{1, 2, 3, 3, 4, 5}
intersection	RDD containing only elements found in both RDDs	<code>rdd.intersection(other)</code>	{3}
subtract	Remove the contents of one RDD (e.g. remove training data)	<code>rdd.subtract(other)</code>	{1, 2}
cartesian	Cartesian product with the other RDD	<code>rdd.cartesian(other)</code>	{(1, 3), (1, 4), ... (3,5)}

As you can see there are a wide variety of transformations available on all RDDs regardless of our specific underlying data. We can transform our data element-wise, obtain distinct elements, and do a variety of set operations.



## Actions

The most common action on basic RDDs you will likely use is `reduce`. `Reduce` takes in a function which operates on two elements of the same type of your RDD and returns a new element of the same type. A simple example of such a function is `+`, which we can use to sum our RDD. With `reduce` we can easily sum the elements of our RDD, count the number of elements, and perform other types of aggregations.

Example 3-29. Python `reduce` example

```
sum = rdd.reduce(lambda x, y: x + y)
```

Example 3-30. Scala `reduce` example

```
val sum = rdd.reduce((x, y) => x + y)
```

Example 3-31. Java `reduce` example

```
Integer sum = rdd.reduce(new Function2<Integer, Integer, Integer>() {  
    public Integer call(Integer x, Integer y) { return x + y;}  
});
```

Similar to `reduce` is `fold` which also takes a function with the same signature as needed for `reduce`, but also takes a “zero value” to be used for the initial call on each partition. The zero value you provide should be the identity element for your operation, that is applying it multiple times with your function should not change the value, (e.g. 0 for `+`, 1 for `*`, or an empty list for concatenation).

### Tip

You can minimize object creation in `fold` by modifying and returning the first of the two parameters in-place. However, you should not modify the second parameter.

`Fold` and `reduce` both require that the return type of our result be the same type as that of the RDD we are operating over. This works well for doing things like sum, but sometimes we want to return a different type. For example when computing the running average we need to have a different return type. We could implement this using a `map` first where we transform every element into the element and the number 1 so that the `reduce` function can work on pairs.

The `aggregate` function frees us from the constraint of having the return the same type as the RDD which we are working on. With `aggregate`, like `fold`, we supply an initial zero value of the type we want to return. We then supply a function to combine the elements from our RDD with the accumulator. Finally, we need to supply a second function to merge two accumulators, given that each node accumulates its own results locally.

We can use aggregate to compute the average of a RDD avoiding a map before the fold.

Example 3-32. Python aggregate example

```
sumCount = nums.aggregate((0, 0),
                           (lambda x, y: (x[0] + y, x[1] + 1),
                            (lambda x, y: (x[0] + y[0], x[1] + y[1]))))
return sumCount[0] / float(sumCount[1])
```

Example 3-33. Scala aggregate example

```
val result = input.aggregate((0, 0))((
    (x, y) => (x._1 + y, x._2 + 1),
    (x, y) => (x._1 + y._1, x._2 + y._2))
val avg = result._1 / result._2.toDouble
```

Example 3-34. Java aggregate example

```
class AvgCount {
    public AvgCount(int total, int num) {
        this.total = total;
        this.num = num;
    }
    public int total;
    public int num;
    public double avg() {
        return total / (double) num;
    }
}

Function2<AvgCount, Integer, AvgCount> addAndCount =
    new Function2<AvgCount, Integer, AvgCount>() {
        public AvgCount call(AvgCount a, Integer x) {
            a.total += x;
            a.num += 1;
            return a;
        }
    };

Function2<AvgCount, AvgCount, AvgCount> combine =
    new Function2<AvgCount, AvgCount, AvgCount>() {
        public AvgCount call(AvgCount a, AvgCount b) {
            a.total += b.total;
            a.num += b.num;
            return a;
        }
    }
```

```
};
AvgCount initial = new AvgCount(o, o);
AvgCount result = rdd.aggregate(initial, addAndCount, combine);
System.out.println(result.avg());
```

Some actions on RDDs return some or all of the data to our driver program in the form of a regular collection or value.

The simplest and most common operation that returns data to our driver program is `collect()`, which returns the entire RDD's contents. `collect` suffers from the restriction that all of your data must fit on a single machine, as it all needs to be copied to the driver.

`take(n)` returns `n` elements from the RDD and attempts to minimize the number of partitions it accesses, so it may represent a biased collection. Its important to note that these operations do not return the elements in the order you might expect.

These operations are useful for unit tests and quick debugging, but may introduce bottlenecks when dealing with large amounts of data.

If there is an ordering defined on our data, we can also extract the top elements from an RDD using `top`. `top` will use the default ordering on the data, but we can supply our own comparison function to extract the top elements.

Sometimes we need a sample of our data in our driver program. The `takeSample(withReplacement, num, seed)` function allows us to take a sample of our data either with or without replacement. For more control, we can create a Sampled RDD and `collect` which we will talk about in the Sampling your data section in the Simple Optimizations chapter.

The further standard operations on a basic RDD all behave pretty much exactly as one would imagine from their name. `count()` returns a count of the elements, and `countByKey()` returns a map of each unique value to its count. See the [basic RDD actions table](#) for more actions.

Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}

Function Name	Purpose	Example	Result
<code>collect()</code>	Return all elements from the RDD	<code>rdd.collect()</code>	{1, 2, 3, 3}
<code>count()</code>	Number of elements in the RDD	<code>rdd.count()</code>	4
<code>take(num)</code>	Return num elements from the RDD	<code>rdd.take(2)</code>	{1, 2}

Function Name	Purpose	Example	Result
<code>top(num)</code>	Return the top num elements the RDD	<code>rdd.top(2)</code>	<code>{3, 3}</code>
<code>takeOrdered(num)(ordering)</code>	Return num elements based on providing ordering	<code>rdd.takeOrdered(2)(myOrdering)</code>	<code>{3, 3}</code>
<code>takeSample(withReplacement, num, [seed])</code>	Return num elements at random	<code>rdd.takeSample(false, 1)</code>	non-deterministic
<code>reduce(func)</code>	Combine the elements of the RDD together in parallel (e.g. sum)	<code>rdd.fold((x, y) =&gt; x + y)</code>	9
<code>fold(zero)(func)</code>	Same as reduce but with the provided zero value	<code>rdd.fold(o)((x, y) =&gt; x + y)</code>	9
<code>aggregate(zeroValue)(seqOp, combOp)</code>	Similar to reduce but used to return a different type	<code>rdd.aggregate(o, o) ({case (x, y) =&gt; (y._1() + x, y._2() + 1)}, {case (x, y) =&gt; (y._1() + x._1(), y._2() + x._2()) })</code>	(9, 4)
<code>foreach(func)</code>	Apply the provided function to each element of the RDD	<code>rdd.foreach(func)</code>	nothing

## Converting Between RDD Types

We don't have to do anything special to get back the correct templated/generic type of RDD (that is, our RDD of Strings can become an RDD of Integers just by calling `map` with the correct function). Some functions are only available on certain types of RDDs, such as `average` on numeric RDDs and `join` on key-value pair RDDs. We will cover these special functions for numeric data in (to come) and pair RDDs in [Chapter 4](#).

In Scala and Java, these methods aren't defined on the standard RDD class, so to access this additional functionality we have to make sure we get the correct specialized class.

## Scala

In Scala the conversion between such RDDs (like from `RDD[Double]` and `RDD[Numeric]` to `DoubleRDD`) is handled automatically using implicit conversions. As mentioned in standard

imports, we need to add `import org.apache.spark.SparkContext._` for these conversion to work. You can see the implicit conversions listed in the object `SparkContext` in the Spark source code at `./core/src/main/scala/org/apache/spark/SparkContext.scala`. These implicits also allow for RDDs of Scala types to be written out to HDFS and similar.

Implicits, while quite powerful, can sometimes be confusing. If you call a function like `stats()` on an RDD, you might look at the [scaladocs for the RDD class](#) and notice there is no `stats()` function. The call manages to succeed because of implicit conversions between `RDD[Numeric]` and `DoubleRDDFunctions`. When looking for functions on your RDD in Scaladoc make sure to look at functions that are available in the other RDD classes.

## Java

In Java the conversion between the specialized types of RDDs is a bit more explicit. This has the benefit of giving you a greater understanding of what exactly is going on, but can be a bit more cumbersome.

To allow Spark to determine the correct return type, instead of always using the `Function` class we will need to use specialized versions. If we want to create a `DoubleRDD` from an RDD of type `T`, rather than using `Function<T, Double>` we use `DoubleFunction<T>`. The [special Java functions table](#) shows the specialized functions and their uses.

We also need to call different functions on our RDD (so we can't just create a `DoubleFunction` and pass it to `map`). When we want a `DoubleRDD` back instead of calling `map` we need to call `mapToDouble` with the same pattern followed with all other functions.

Table 3-5. Java interfaces for type-specific functions

Function name	Equivalent Function* <code>&lt;A, B,...&gt;</code>	Usage
<code>DoubleFlatMapFunction&lt;T&gt;</code>	<code>Function&lt;T, Iterable&lt;Double&gt;&gt;</code>	<code>DoubleRDD</code> from a <code>flatMapToDouble</code>
<code>DoubleFunction&lt;T&gt;</code>	<code>Function&lt;T, double&gt;</code>	<code>DoubleRDD</code> from <code>mapToDouble</code>
<code>PairFlatMapFunction&lt;T, K, V&gt;</code>	<code>Function&lt;T, Iterable&lt;Tuple2&lt;K, V&gt;&gt;&gt;</code>	<code>PairRDD&lt;K, V&gt;</code> from a <code>flatMapToPair</code>
<code>PairFunction&lt;T, K, V&gt;</code>	<code>Function&lt;T, Tuple2&lt;K, V&gt;&gt;</code>	<code>PairRDD&lt;K, V&gt;</code> from a <code>mapToPair</code>

We can modify our previous example where we squared an RDD of numbers to produce a `JavaDoubleRDD`. This gives us access to the additional `DoubleRDD` specific functions like `average` and `stats`.

Example 3-35. Java create `DoubleRDD` example

```
JavaDoubleRDD result = rdd.mapToDouble(  
    new DoubleFunction<Integer>() {  
        public double call(Integer x) {  
            return (double) x * x;  
        }  
    });  
System.out.println(result.average());
```

## Python

The Python API is structured a bit different than both the Java and Scala API. Like the Scala API, we don't need to be explicit to access the functions which are only available on `Double` or `Pair` RDDs. In Python all of the functions are implemented on the base `RDD` class and will simply fail at runtime if the type doesn't work.

## Persistence (Caching)

As discussed earlier, Spark RDDs are lazily evaluated, and sometimes we may wish to use the same RDD multiple times. If we do this naively, Spark will recompute the RDD and all of its dependencies each time we call an action on the RDD. This can be especially expensive for iterative algorithms, which look at the data many times. Another trivial example would be doing a count and then writing out the same RDD.

Example 3-36. Scala double execute example

```
val result = input.map(x => x*x)  
println(result.count())  
println(result.collect().mkString(","))
```

To avoid computing an RDD multiple times, we can ask Spark to *persist* the data. When we ask Spark to persist an RDD, the nodes that compute the RDD store their partitions. If a node that has data persisted on it fails, Spark will recompute the lost partitions of the data when needed. We can also replicate our data on multiple nodes if we want to be able to handle node failure without slowdown.

Spark has [many levels of persistence](#) to choose from based on what our goals are. In Scala and Java, the default `persist()` will store the data in the JVM heap as unserialized objects. In Python, we always serialize the data that `persist` stores, so the default is instead stored in the JVM heap as pickled objects. When we write data out to disk or off-heap storage that data is also always serialized.

## Tip

Off-heap caching is experimental and uses [Tachyon](#). If you are interested in off-heap caching with Spark, [take a look at the running Spark on Tachyon guide](#).

Table 3-6. Persistence levels

Level	Space Used	CPU time	In memory	On Disk	Nodes with data	Comments
MEMORY_ONLY	High	Low	Y	N	1	
MEMORY_ONLY_2	High	Low	Y	N	2	
MEMORY_ONLY_SER	Low	High	Y	N	1	
MEMORY_ONLY_SER_2	Low	High	Y	N	2	
MEMORY_AND_DISK	High	Medium	Some	Some	1	Spills to disk if there is too much data to fit in memory.
MEMORY_AND_DISK_2	High	Medium	Some	Some	2	Spills to disk if there is too much data to fit in memory.
MEMORY_AND_DISK_SER	Low	High	Some	Some	1	Spills to disk if there is too much data to fit in memory.
MEMORY_AND_DISK_SER_2	Low	High	Some	Some	2	Spills to disk if there is too much data to fit in memory.
DISK_ONLY	Low	High	N	Y	1	
DISK_ONLY_2	Low	High	N	Y	2	

Example 3-37. Scala `persist` example

```
val result = input.map(x => x*x)
result.persist(MEMORY_ONLY)
println(result.count())
println(result.collect().mkString(","))
```

## Tip

You will note that we called `persist` on the RDD before the first action. The `persist` call on its own doesn't force evaluation.

If you attempt to cache too much data to fit in memory, Spark will automatically evict old partitions using a Least Recently Used (LRU) cache policy. For the memory-only storage levels, it will recompute these partitions the next time they are accessed, while for the memory-and-disk ones, it will write them out to disk. In either case, this means that you don't have to worry about your job breaking if you ask Spark to cache too much data. However, caching unnecessary data can lead to eviction of useful data and more recomputation time.

Finally, RDDs come with a method called `unpersist()` that lets you manually remove them from the cache.

## Conclusion

In this chapter, we have covered the RDD execution model and a large number of common operations on RDDs. If you have gotten here, congratulations—you've learned all the core concepts of working in Spark. In the next chapter, we'll cover a special set of operations available on RDDs of key-value pairs, which are the most common way to aggregate or group together data in parallel. After that, we discuss input and output from a variety of data sources, and more advanced topics in working with `SparkContext`.

---

[8] The ability to always recompute an RDD is actually why RDDs are called “resilient”. When a machine holding RDD data fails, Spark uses this ability to recompute the missing partitions, transparent to the user.



## Chapter 4. Working with Key-Value Pairs

This chapter covers how to work with RDDs of key-value pairs, which are a common data type required for many operations in Spark. Key-value RDDs expose new operations such as aggregating data items by key (e.g., counting up reviews for each product), grouping together data with the same key, and grouping together two different RDDs. Oftentimes, to work with data records in Spark, you will need to turn them into key-value pairs and apply one of these operations.

We also discuss an advanced feature that lets users control the layout of pair RDDs across nodes: *partitioning*. Using controllable partitioning, applications can sometimes greatly reduce communication costs, by ensuring that data that will be accessed together will be on the same node. This can provide significant speedups. We illustrate partitioning using the PageRank algorithm as an example. Choosing the right partitioning for a distributed dataset is similar to choosing the right data structure for a local one—in both cases, data layout can greatly affect performance.

### Motivation

Spark provides special operations on RDDs containing key-value pairs. These RDDs are called Pair RDDs. Pair RDDs are a useful building block in many programs, as they expose operations that allow you to act on each key in parallel or regroup data across the network. For example, pair RDDs have a `reduceByKey` method that can aggregate data separately for each key, and a `join` method that can merge two RDDs together by grouping elements with the same key. It is common to extract fields from an RDD (representing for instance, an event time, customer ID, or other identifier) and use those fields as keys in Pair RDD operations.

### Creating Pair RDDs

There are a number of ways to get Pair RDDs in Spark. Many formats we explore in [Chapter 5](#) will directly return Pair RDDs for their key-value data. In other cases we have a regular RDD that we want to turn into a Pair RDDs. To illustrate creating a Pair RDDs we will key our data by the first word in each line of the input.

In Python, for the functions on keyed data to work we need to make sure our RDD consists of tuples.

Example 4-1. Python create pair RDD using the first word as the key

```
input.map(lambda x: (x.split(" ")[0], x))
```

In Scala, to create Pair RDDs from a regular RDD, we simply need to return a tuple from our function.

Example 4-2. Scala create pair RDD using the first word as the key

```
input.map(x => (x.split(" ")(0), x))
```

Java doesn't have a built-in tuple type, so Spark's Java API has users create tuples using the `scala.Tuple2` class. This class is very simple: Java users can construct a new tuple by writing `new Tuple2(elem1, elem2)` and can then access the elements with the `._1()` and `._2()` methods.

Java users also need to call special versions of Spark's functions when creating Pair RDDs. For instance, the `mapToPair` function should be used in place of the basic `map` function. This is discussed in more detail in [converting between RDD types](#), but let's look at a simple example below.

Example 4-3. Java create pair RDD using the first word as the key

```
PairFunction<String, String, String> keyData =  
    new PairFunction<String, String, String>() {  
        public Tuple2<String, String> call(String x) {  
            return new Tuple2(x.split(" ")[0], x);  
        }  
    };  
JavaPairRDD<String, String> rdd = input.mapToPair(keyData);
```

When creating a Pair RDD from an in memory collection in Scala and Python we only need to make sure the types of our data are correct, and call `parallelize`. To create a Pair RDD in Java from an in memory collection we need to make sure our collection consists of tuples and also call `SparkContext.parallelizePairs` instead of `SparkContext.parallelize`.

## Transformations on Pair RDDs

Pair RDDs are allowed to use all the transformations available to standard RDDs. The same rules from [passing functions to spark](#) apply. Since Pair RDDs contain tuples, we need to pass functions that operate on tuples rather than on individual elements.

In Java and Scala when we run a `map` or `filter` or similar over a Pair RDDs, our function will be passed an instance of `scala.Tuple2`. In Scala pattern matching is a common way of extracting the

individual values, whereas in Java we use the `._1()` and `._2()` values to access the elements. In Python our Pair RDDs consist of standard Python tuple objects that we interact with as normal.

For instance, we can create take our Pair RDD from the previous section then filter out lines longer than 20 characters.

Example 4-4. Python simple filter on second element

```
result = pair.filter(lambda x: len(x[1]) < 20)
```

Example 4-5. Scala simple filter on second element

```
pair.filter{case (x, y) => y.length < 20}
```

Example 4-6. Java simple filter on second element

```
Function<Tuple2<String, String>, Boolean> longWordFilter =  
    new Function<Tuple2<String, String>, Boolean>() {  
        public Boolean call(Tuple2<String, String> input) {  
            return (input._2().length() < 20);  
        }  
    };  
JavaPairRDD<String, String> result = rdd.filter(longWordFilter);
```

Sometimes working with these pairs can be awkward if we only want to access the value part of our Pair RDD. Since this a common pattern, Spark provides the `mapValues(func)` function which is the same as `map{case (x, y) => (x, func(y))}` and we will use this function in many of our examples.

## Aggregations

When datasets are described in terms of key-value pairs, it is common to want to aggregate statistics across all elements with the same key. We have looked at the fold, combine, and reduce actions on basic RDDs, and similar per-key transformations exist on Pair RDDs. Spark has a similar set of operations which combine the values together which have the same key. Instead of being actions these operations return RDDs and as such are transformations.

`reduceByKey` is quite similar to `reduce`, both take a function and use it to combine values. `reduceByKey` runs several parallel reduce operations, one for each key in the dataset, where each operation combines values together which have the same key. Because datasets can have very large numbers of keys, `reduceByKey` is not implemented as an action that returns a value back to the user program. Instead, it returns a new RDD consisting of each key and the reduced value for that key.

foldByKey is quite similar to fold, both use a zero value of the same type of the data in our RDD and combination function. Like with fold the provided zero value for foldByKey should have no impact when added with your combination function to another element.

We can use reduceByKey along with mapValues to compute the per-key average in a very similar manner to how we used fold and map compute the entire RDD average. As with averaging, we can achieve the same result using a more specialized function we will cover next.

Example 4-7. Python per key average with reduceByKey and mapValues

```
rdd.mapValues(lambda x: (x, 1)).reduceByKey(lambda x, y: (x[0] + y[0], x[1] + y[1]))
```

Example 4-8. Scala per key average with reduceByKey and mapValues

```
rdd.mapValues(x => (x, 1)).reduceByKey((x, y) => (x._1 + y._1, x._2 + y._2))
```

## Note

Those familiar with the combiner concept from MapReduce should note that calling reduceByKey and foldByKey will automatically perform combining locally on each machine before computing global totals for each key. The user does not need to specify a combiner. The more general combineByKey interface allows you to customize combining behavior.

We can use a similar approach to also implement the classic distributed word count problem. We will use flatMap from the previous chapter so that we can produce a Pair RDD of words and the number 1 and then sum together all of the words using reduceByKey like in our previous example.

Example 4-9. Python word count example

```
rdd = sc.textFile("s3://...")
words = rdd.flatMap(lambda x: x.split(" "))
result = words.map(lambda x: (x, 1)).reduceByKey(lambda x, y: x + y)
```

Example 4-10. Scala word count example

```
val input = sc.textFile("s3://...")
val words = input.flatMap(x => x.split(" "))
val result = words.map(x => (x, 1)).reduceByKey((x, y) => x + y)
```

Example 4-11. Java word count example

```
JavaRDD<String> input = sc.textFile("s3://...")
JavaRDD<String> words = rdd.flatMap(new FlatMapFunction<String, String>() {
    public Iterable<String> call(String x) { return Arrays.asList(x.split(" ")); }
});
JavaPairRDD<String, Integer> result = words.mapToPair(
    new PairFunction<String, String, Integer>() {
```

```

    public Tuple2<String, Integer> call(String x) { return new Tuple2(x, 1); }
  }).reduceByKey(
    new Function2<Integer, Integer, Integer>() {
      public Integer call(Integer a, Integer b) { return a + b; }
    });

```

## Tip

We can actually implement word count even faster by using the `countByValue()` function on the first RDD: `lines.flatMap(x => x.split(" ")).countByValue()`.

`combineByKey` is the most general of the per-key aggregation functions and provides flexibility in how the values associated with each key are combined. Most of the other per-key combiners are implemented using it. Like `aggregate`, `combineByKey` allows the user to return values which are not the same type as our input data. To use `combineByKey` we need to provide a number of different functions.

The first required function is called `createCombiner` which should take a single element in the source RDD and return an element of the desired type in the resulting RDD. A good example of this would be if we were building a list of all of the values for each key, `createCombiner` could return a list containing the element it was passed in. In implementing `foldByKey` `createCombiner` creates a new instance of the provided zero value and combines it with the input value, and in `reduceByKey` it is the identity operator (namely it just returns the input).

The second required function is `mergeValue` which takes the current accumulated value for the key and the new value and returns a new accumulated value for the key. If we wanted to make a list of elements we might have `mergeValue` simply append the new element to the current list. With `reduceByKey` and `foldByKey` the `mergeValue` function is used is simply the user provided merge function. `mergeValue` is used to update the accumulated value for each key when processing a new element.

The final required function you need to provide to `combineByKey` is `mergeCombiners`. Since we don't run through the elements linearly, we can have multiple accumulators for each key. `mergeCombiners` must take two accumulators (of the type returned by `createCombiner`) and return a merged result. If we were using `combineByKey` to implement `groupByKey` our `mergeCombiners` function could just return the lists as appended to each other. In the case of `foldByKey` and `reduceByKey` since our accumulator is the same type as our data, the `combineByKey` function they use is the same as the `mergeValue` function.



```

        return a;
    }
};
AvgCount initial = new AvgCount(0,0);
JavaPairRDD<String, AvgCount> avgCounts =
    rdd.combineByKey(createAcc, addAndCount, combine);
Map<String, AvgCount> countMap = avgCounts.collectAsMap();
for (Entry<String, AvgCount> entry : countMap.entrySet()) {
    System.out.println(entry.getKey() + ":" + entry.getValue().avg());
}

```

There are many options for combining our data together by key. Most of them are implemented on top of `combineByKey` but provide a simpler interface. Using one of the specialized per-key combiners in Spark can be much faster than the naive approach of grouping our data and then reducing the data.

## Tuning the Level of Parallelism

So far we have talked about how all of our transformations are distributed, but we have not really looked at how Spark decides how to split up the work. Every RDD has a fixed number of *partitions* which determine the degree of parallelism to use when executing operations on the RDD.

When performing aggregations or grouping operations, we can ask Spark to use a specific number of partitions. Spark will always try to infer a sensible default value based on the size of your cluster, but in some cases users will want to tune the level of parallelism for better performance.

Most of the operators discussed in this chapter accept a second parameter giving the number of partitions to use when creating the grouped or aggregated RDD:

Example 4-15. Python `reduceByKey` with custom parallelism

```

data = [("a", 3), ("b", 4), ("a", 1)]
sc.parallelize(data).reduceByKey(lambda x, y: x + y)      # Default parallelism
sc.parallelize(data).reduceByKey(lambda x, y: x + y, 10) # Custom parallelism

```

Example 4-16. Scala `reduceByKey` with custom parallelism

```

val data = Seq(("a", 3), ("b", 4), ("a", 1))
sc.parallelize(data).reduceByKey(_ + _)      // Default parallelism
sc.parallelize(data).reduceByKey(_ + _, 10) // Custom parallelism

```

Sometimes, we want to change the partitioning of an RDD outside of the context of grouping and aggregation operations. For those cases, Spark provides the repartition function. Keep in mind that repartitioning your data is a fairly expensive operation. Spark also has an optimized version of repartition called coalesce that allows avoiding data movement, but only if you are decreasing the number of RDD partitions. To know whether you can safely call coalesce you can check the size of the RDD using `rdd.partitions.size()` in Java/Scala and `rdd.getNumPartitions()` in Python and make sure that you are coalescing it fewer partitions that it currently has.

## Grouping Data

With keyed data a common use case is grouping our data together by key, say joining all of a customers orders together.

If our data is already keyed in the way which we are interested `groupByKey` will group our data together using the key in our RDD. On an RDD consisting of keys of type `K` and values of type `V` we get back an RDD of type `[K, Iterable[V]]`.

`groupBy` works on unpaired data or data where we want to use a different condition besides equality on the current key. `groupBy` takes a function which it applies to every element in the source RDD and uses the result to determine the key.

## Tip

If you find your self writing code where you `groupByKey` and then do a reduce or fold on the values, you can probably achieve the same result more efficiently by using one of the per-key combiners. Rather than reducing the RDD down to an in memory value, the data is reduced per-key and we get back an RDD with the reduced values corresponding to each key. e.g. `rdd.reduceByKey(func)` produces the same RDD as `rdd.groupByKey().mapValues(value => value.reduce(func))` but is more efficient as it avoids the step of creating a list of values for each key.

In addition to grouping together data from a single RDD, we can group together data sharing the same key from multiple RDDs using a function called `cogroup` `cogroup` over two RDDs sharing the same key type `K` with the respective value types `V` and `W` gives use back `RDD[(K, Tuple(Iterable[V], Iterable[W]))]`. If one of the RDDs doesn't have an elements for a given key that is present in the other RDD the corresponding Iterable is simply empty. `cogroup` gives us the power to group together data from multiple RDDs.

The basic transformation on which the joins we discuss in the next section are implemented is `cogroup`. `cogroup` returns a Pair RDD with an entry for every key found in any of the RDDs it is



called on along with a tuple of iterators with each iterator containing all of the elements in the corresponding RDD for that key.

## Tip

cogroup can be used for much more than just implementing joins. We can also use it to implement intersect by key. Additionally, cogroup can work on three RDDs at once.

## Joins

Some of the most useful operations we get with keyed data comes from using it together with other keyed data. Joining data together is probably one of the most common operations on a Pair RDD, and we have a full range of options including right and left outer joins, cross joins, inner joins.

The simple join operator is an inner join. Only keys which are present in both Pair RDDs are output. When there are multiple values for the same key in one of the inputs the resulting Pair RDD will also have multiple entries for the same key, with the values being the Cartesian product of the values for that key in each of the input RDDs. A simple way to understand this is by looking at an example of a join.

Example 4-17. Scala shell inner join example

```
storeAddress = {  
  (Store("Ritual"), "1026 Valencia St"), (Store("Philz"), "748 Van Ness Ave"),  
  (Store("Philz"), "3101 24th St"), (Store("Starbucks"), "Seattle")}
```

```
storeRating = {  
  (Store("Ritual"), 4.9), (Store("Philz"), 4.8)}
```

```
storeAddress.join(storeRating) = {  
  (Store("Ritual"), ("1026 Valencia St", 4.9)),  
  (Store("Philz"), ("748 Van Ness Ave", 4.8)),  
  (Store("Philz"), ("3101 24th St", 4.8))}
```

Sometimes we don't need the key to be present in both RDDs to want it in our result. For example if we were joining customer information with recommendations we might not want to drop customers if there were not any recommendations yet. `leftOuterJoin(other)` and `rightOuterJoin(other)` both join Pair RDDs together by key where one of the Pair RDDs can be missing the key.

With `leftOuterJoin` the resulting Pair RDD has entries for each key in the source RDD. The value associated with each key in the result is a tuple of the value from the source RDD and an Option (or Optional in Java) for the value from the other Pair RDD. In Python if a value isn't present `None` is used and if the value is present the regular value, without any wrapper, is used. Like with `join` we can have multiple entries for each key and when this occurs we get the cartesian product between the two list of values.

## Tip

Optional is part of [Google's Guava library](#) and is similar to nullable. We can check `isPresent()` to see if its set, and `get()` will return the contained instance provided it has data present.

`rightOuterJoin` is almost identical to `leftOuterJoin` except the key must be present in the other RDD and the tuple has an option for the source rather than other RDD.

We can look at our example from earlier and do a `leftOuterJoin` and a `rightOuterJoin` between the two Pair RDDs we used to illustrate join.

Example 4-18. Scala shell `leftOuterJoin` / `rightOuterJoin` examples

```
storeAddress.leftOuterJoin(storeRating) =  
{(Store("Ritual"),("1026 Valencia St",Some(4.9))),  
  (Store("Starbucks"),("Seattle",None)),  
  (Store("Philz"),("748 Van Ness Ave",Some(4.8))),  
  (Store("Philz"),("3101 24th St",Some(4.8)))}
```

```
storeAddress.rightOuterJoin(storeRating) =  
{(Store("Ritual"),(Some("1026 Valencia St"),4.9)),  
  (Store("Philz"),(Some("748 Van Ness Ave"),4.8)),  
  (Store("Philz"), (Some("3101 24th St"),4.8))}
```

## Sorting Data

Having sorted data is quite useful in many cases, especially when producing downstream output. We can sort an RDD with key value pairs provided that there is an ordering defined on the key. Once we have sorted our data any subsequent call on the sorted data to collect or save will result in ordered data.

Since we often want our RDDs in the reverse order, the `sortByKey` function takes a parameter called `ascending` indicating if we want it in ascending order (defaults to true). Sometimes we want a different sort order entirely, and to support this we can provide our own comparison

function here we will sort our RDD by converting the integers to strings and using the string comparison functions.

Example 4-19. Custom sort order in Python sorting integers as if strings

```
rdd.sortByKey(ascending=True, numPartitions=None, keyfunc = lambda x: str(x))
```

Example 4-20. Custom sort order in Scala sorting integers as if strings

```
val input: RDD[(Int, Venue)] = ...
```

```
implicit val sortIntegersByString = new Ordering[Int] {
```

```
  override def compare(a: Int, b: Int) = a.toString.compare(b.toString)
```

```
}
```

```
rdd.sortByKey()
```

Example 4-21. Custom sort order in Java sorting integers as if strings

```
class IntegerComparator implements Comparator<Integer> {
```

```
  public int compare(Integer a, Integer b) {
```

```
    return String.valueOf(a).compareTo(String.valueOf(b))
```

```
  }
```

```
}
```

```
rdd.sortByKey(comp)
```

The following tables summarize transformations on pair RDDs.

Table 4-1. Transformations on one Pair RDD (example ({(1, 2), (3, 4), (3, 6)}))

Function Name	Purpose	Example	Result
combineByKey(createCombiner, mergeValue, mergeCombiners, partitioner)	Combine values with the same key together	See <a href="#">combine by key example</a>	
groupByKey()	Group together values with the same key	rdd.groupByKey()	{{(1, [2]), (3, [4, 6])}}
reduceByKey(func)	Combine values with the same key together	rdd.reduceByKey( (x, y) => x + y)	{{(1, 2), (3, 10)}}
mapValues(func)	Apply a function to each value of a Pair RDD without changing the key	rdd.mapValues(x => x+1)	{{(1, 3), (3, 5), (3, 7)}}
flatMapValues(func)	Apply a function which returns an iterator to each value of a Pair RDD and for each element returned produce a key-value entry	rdd.flatMapValues(x => x.to(5))	{{(1,2), (1,3), (1,4), (1,5), (3, 4),

Function Name	Purpose	Example	Result
	with the old key. Often used for tokenization.		(3,5)}
keys()	Return an RDD of just the keys	rdd.keys()	{1, 3, 3}
values()	Return an RDD of just the values	rdd.values()	{2, 4, 6}
sortByKey()	Returns an RDD sorted by the key	rdd.sortByKey()	{(1, 2), (3, 4), (3, 6)}

Table 4-2. Transformations on two Pair RDD (example {(1, 2), (3, 4), (3, 6)}) other {(3, 9)}

Function Name	Purpose	Example	Result
subtractByKey	Remove elements with a key present in the other RDD	rdd.subtractByKey(other)	{1, 2}
join	Perform an inner join between two RDDs	rdd.join(other)	{(3, (4, 9)), (3, (6, 9))}
rightOuterJoin	Perform a join between two RDDs where the key must be present in the first RDD.	rdd.rightOuterJoin(other)	{(3,(Some(4),9)), (3,(Some(6),9))}
leftOuterJoin	Perform a join between two RDDs where the key must be present in the other RDD.	rdd.leftOuterJoin(other)	{(1,(2,None)), (3,(4,Some(9))), (3,(6,Some(9)))}
cogroup	Group together data from both RDDs sharing the same key	rdd.cogroup(other)	{(1,([2,[]]), (3,([4, 6],[9]))}

## Actions Available on Pair RDDs

Like with the transformations, all of the traditional actions available on the base RDD are also available on Pair RDDs. Some additional actions are available on Pair RDDs which take advantage of the key-value nature of the data.

Table 4-3. Actions on Pair RDDs (example  $\{(1, 2), (3, 4), (3, 6)\}$ )

<code>countByKey()</code>	Count the number of elements for each key	<code>rdd.countByKey()</code>	$\{(1, 1), (3, 2)\}$
<code>collectAsMap()</code>	Collect the result as a map to provide easy lookup	<code>rdd.collectAsMap()</code>	$\text{Map}\{(1, 2), (3, 4), (3, 6)\}$
<code>lookup(key)</code>	Return all values associated with the provided key	<code>rdd.lookup(3)</code>	$[4, 6]$

There are also multiple other actions on Pair RDDs that save the RDD, which we will examine in [the next chapter](#).

## Data Partitioning

The final Spark feature we will discuss in this chapter is how to control datasets' partitioning across nodes. In a distributed program, communication is very expensive, so laying out data to minimize network traffic can greatly improve performance. Much like how a single-node program needs to choose the right data structure for a collection of records, Spark programs can choose to control their RDDs' partitioning to reduce communication. Partitioning will not be helpful in all applications—for example, if a given RDD is only scanned once, there is no point in partitioning it in advance. It is only useful when a dataset is reused *multiple times* in key-oriented operations such as joins. We will give some examples below.

Spark's partitioning is available on all RDDs of key-value pairs, and causes the system to group together elements based on a function of each key. Although Spark does not give explicit control of which worker node each key goes to (partly because the system is designed to work even if specific nodes fail), it lets the program ensure that a *set* of keys will appear together on *some* node. For example, one might choose to hash-partition an RDD into 100 partitions so that keys that have the same hash value modulo 100 appear on the same node. Or one might range-partition the RDD into sorted ranges of keys so that elements with keys in the same range appear on the same node.

As a simple example, consider an application that keeps a large table of user information in memory—say, an RDD of (UserID, UserInfo) pairs where UserInfo contains a list of topics the user is subscribed to. The application periodically combines this table with a smaller file representing events that happened in the past five minutes—say, a table of (UserID, LinkInfo) pairs for users who have clicked a link on a website in those five minutes. For example, we may wish to count how many users visited a link that was *not* to one of their subscribed topics. We can perform this combining with Spark's join operation, which can be used to group the UserInfo and LinkInfo pairs for each UserID by key. Our application would look like this:

```
// Initialization code; we load the user info from a Hadoop SequenceFile on HDFS.
// This distributes elements of userData by the HDFS block where they are found,
// and doesn't provide Spark with any way of knowing in which partition a
// particular UserID is located.
val sc = new SparkContext(...)
val userData = sc.sequenceFile[UserID, UserInfo]("hdfs://...").persist()

// Function called periodically to process a log file of events in the past 5 minutes;
// we assume that this is a SequenceFile containing (UserID, LinkInfo) pairs.
def processNewLogs(logFileName: String) {
  val events = sc.sequenceFile[UserID, LinkInfo](logFileName)
  val joined = userData.join(events) // RDD of (UserID, (UserInfo, LinkInfo)) pairs
  val offTopicVisits = joined.filter {
    case (userId, (userInfo, linkInfo)) => // Expand the tuple into its components
      !userInfo.topics.contains(linkInfo.topic)
  }.count()
  println("Number of visits to non-subscribed topics: " + offTopicVisits)
}
```

This code will run fine as is, but it will be inefficient. This is because the join operation, called each time that processNewLogs is invoked, does not know anything about how the keys are partitioned in the datasets. By default, this operation will hash all the keys of *both* datasets, sending elements with the same key hash across the network to the same machine, and then join on that machine the elements with the same key (to come). Because we expect the userData table to be much larger than the small log of events seen every five minutes, this wastes a lot of work: the userData table is hashed and shuffled across the network on every call, even though it doesn't change.

Fixing this is simple: just use the partitionBy transformation on userData to hash-partition it at the start of the program. We do this by passing a spark.HashPartitioner object to partitionBy:

```
val sc = new SparkContext(...)
val userData = sc.sequenceFile[UserID, UserInfo]("hdfs://...")
    .partitionBy(new HashPartitioner(100)) // Create 100 partitions
    .persist()
```

The processNewLogs method can remain unchanged—the events RDD is local to processNewLogs, and is only used once within this method, so there is no advantage in specifying a partitioner for events. Because we called partitionBy when building userData, Spark will now know that it is hash-partitioned, and calls to join on it will take advantage of this

information. In particular, when we call `userData.join(events)`, Spark will only shuffle the events RDD, sending events with each particular UserID to the machine that contains the corresponding hash partition of `userData` (to come). The result is that a lot less data is communicated over the network, and the program runs significantly faster.

Note that `partitionBy` is a transformation, so it always returns a *new* RDD—it does not change the original RDD in-place. RDDs can never be modified once created. Therefore it is important to persist and save as `userData` the result of `partitionBy`, not the original `sequenceFile`. Also, the 100 passed to `partitionBy` represents the number of partitions, which will control how many parallel tasks perform further operations on the RDD (e.g., joins); in general, make this at least as large as the number of cores in your cluster.

## Warning

Failure to persist an RDD after it has been transformed with `partitionBy` will cause subsequent uses of the RDD to repeat the partitioning of the data. Without persistence, use of the partitioned RDD will cause re-evaluation of the RDDs complete lineage. That would negate the advantage of `partitionBy`, resulting in repeated partitioning and shuffling of data across the network, similar to what occurs without any specified partitioner.

## Note

When using a `HashPartitioner`, specify a number of hash buckets at least as large as the number of cores in your cluster in order to achieve appropriate parallelism.

In fact, many other Spark operations automatically result in an RDD with known partitioning information; and many operations other than `join` will take advantage of this information. For example, `sortByKey` and `groupByKey` will result in range-partitioned and hash-partitioned RDDs, respectively. On the other hand, operations like `map` cause the new RDD to *forget* the parent's partitioning information, because such operations could theoretically modify the key of each record. The next few sections describe how to determine how an RDD is partitioned, and exactly how partitioning affects the various Spark operations.

## Partitioning in Java and Python

Spark's Java and Python APIs benefit from partitioning the same way as the Scala API. However, in Python, you cannot pass a `HashPartitioner` object to `partitionBy`; instead, you just pass the number of partitions desired (e.g., `rdd.partitionBy(100)`).

## Determining an RDD's Partitioner

In Scala and Java, you can determine how an RDD is partitioned using its partitioner property (or `partitioner()` method in Java).<sup>[9]</sup> This returns a `scala.Option` object, which is a Scala class for a container that may or may not contain one item. (It is considered good practice in Scala to use `Option` for fields that may not be present, instead of setting a field to null if a value is not present, running the risk of a null-pointer exception if the program subsequently tries to use the null as if it were an actual, present value.) You can call `isDefined()` on the `Option` to check whether it has a value, and `get()` to get this value. If present, the value will be a `spark.Partitioner` object. This is essentially a function telling the RDD which partition each key goes into—more about this later.

The partitioner property is a great way to test in the Spark shell how different Spark operations affect partitioning, and to check that the operations you want to do in your program will yield the right result. For example:

```
scala> val pairs = sc.parallelize(List((1, 1), (2, 2), (3, 3)))
pairs: spark.RDD[(Int, Int)] = ParallelCollectionRDD[0] at parallelize at <console>:12
```

```
scala> pairs.partitioner
reso: Option[spark.Partitioner] = None
```

```
scala> val partitioned = pairs.partitionBy(new spark.HashPartitioner(2))
partitioned: spark.RDD[(Int, Int)] = ShuffledRDD[1] at partitionBy at <console>:14
```

```
scala> partitioned.partitioner
res1: Option[spark.Partitioner] = Some(spark.HashPartitioner@5147788d)
```

In this short session, we created an RDD of `(Int, Int)` pairs, which initially have no partitioning information (an `Option` with value `None`). We then created a second RDD by hash-partitioning the first. If we actually wanted to use `partitioned` in further operations, then we should have appended `.cache()` to the third line of input, in which `partitioned` is defined. This is for the same reason that we needed `cache` for `userData` in the previous example: without `cache`, subsequent RDD actions will evaluate the entire lineage of `partitioned`, which will cause `pairs` to be hash-partitioned over and over.

## Operations that Benefit from Partitioning

Many of Spark's operations involve shuffling data by key across the network. All of these will benefit from partitioning. As of Spark 1.0, the operations that benefit from partitioning are:



cogroup, groupWith, join, leftOuterJoin, rightOuterJoin, groupByKey, reduceByKey, combineByKey, and lookup.

For operations that act on a single RDD, such as reduceByKey, running on a pre-partitioned RDD will cause all the values for each key to be computed *locally* on a single machine, requiring only the final, locally-reduced value to be sent from each worker node back to the master. For binary operations, such as cogroup and join, pre-partitioning will cause at least one of the RDDs (the one with the known partitioner) to not be shuffled. If both RDDs have the *same* partitioner, and if they are cached on the same machines (e.g. one was created using mapValues on the other, which preserves keys and partitioning) or if one of them has not yet been computed, then no shuffling across the network will occur.

## Operations that Affect Partitioning

Spark knows internally how each of its operations affects partitioning, and automatically sets the partitioner on RDDs created by operations that partition the data. For example, suppose you called join to join two RDDs; because the elements with the same key have been hashed to the same machine, Spark knows that the result is hash-partitioned, and operations like reduceByKey on the join result are going to be significantly faster.

The flip-side, however, is that for transformations that *cannot* be guaranteed to produce a known partitioning, the output RDD will not have a partitioner set. For example, if you call map on a hash-partitioned RDD of key-value pairs, the function passed to map can in theory change the key of each element, so the result will not have a partitioner. Spark does not analyze your functions to check whether they retain the key. Instead, it provides two other operations, mapValues and flatMapValues, which guarantee that each tuple's key remains the same.

All that said, here are all the operations that result in a partitioner being set on the output RDD: cogroup, groupWith, join, leftOuterJoin, rightOuterJoin, groupByKey, reduceByKey, combineByKey, partitionBy, sort, mapValues (if the parent RDD has a partitioner), flatMapValues (if parent has a partitioner), and filter (if parent has a partitioner). All *other* operations will produce a result with no partitioner.

That there is a partitioner does not answer the question of *what* the output partitioner is for binary operations such as join. By default, it is a hash partitioner, with the number of partitions set to the level of parallelism of the operation. However, if one of the parents has a partitioner object, it will be that partitioner; and if both parents have a partitioner set, it will be the partitioner of the first parent (the one that join was called on, for example).

## Example: PageRank

As an example of a more involved algorithm that can benefit from RDD partitioning, we consider PageRank. The PageRank algorithm, named after Google's Larry Page, aims to assign a measure of importance (a "rank") to each document in a set based on how many documents have links to it. It can be used to rank web pages, of course, but also scientific articles, or influential users in a social network (by treating each user as a "document" and each friend relationship as a "link").

PageRank is an iterative algorithm that performs many joins, so it is a good use case for RDD partitioning. The algorithm maintains two datasets: one of (pageID, linkList) elements containing the list of neighbors of each page, and one of (pageID, rank) elements containing the current rank for each page. It proceeds as follows:

1. Initialize each page's rank to 1.0
2. On each iteration, have page  $p$  send a contribution of  $\text{rank}(p) / \text{numNeighbors}(p)$  to its neighbors (the pages it has links to).
3. Set each page's rank to  $0.15 + 0.85 * \text{contributionsReceived}$ .

The last two steps repeat for several iterations, during which the algorithm will converge to the correct PageRank value for each page. As a simple solution, it's typically enough to run about ten iterations and declare the resulting ranks to be the PageRank values.

Here is the code to implement PageRank in Spark:

```
val sc = new SparkContext(...)

// Assume that our neighbor list was saved as a Spark objectFile
val links = sc.objectFile[(String, Seq[String])]("links")
               .partitionBy(new HashPartitioner(100))
               .persist()

// Initialize each page's rank to 1.0; since we use mapValues, the resulting RDD
// will have the same partitioner as links
var ranks = links.mapValues(_ => 1.0)

// Run 10 iterations of PageRank
for (i <- 0 until 10) {
  val contributions = links.join(ranks).flatMap {
    case (pageId, (links, rank)) =>
```

```

    links.map(dest => (dest, rank / links.size))
  }
  ranks = contributions.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
}

// Write out the final ranks
ranks.saveAsTextFile("ranks")

```

That's it! The algorithm starts with a `ranks` RDD initialized at 1.0 for each element, and keeps updating the `ranks` variable on each iteration. The body of PageRank is pretty simple to express in Spark: it first does a join between the current `ranks` RDD and the static `links` one, in order to obtain the link list and rank for each page ID together, then uses this in a `flatMap` to create “contribution” values to send to each of the page's neighbors. We then add up these values by page ID (i.e. by the page receiving the contribution) and set that page's rank to  $0.15 + 0.85 * \text{contributionsReceived}$ .

Although the code itself is simple, the example does several things to ensure that the RDDs are partitioned in an efficient way, and to minimize communication:

1. Notice that the `links` RDD is joined against `ranks` on each iteration. Since `links` is a static dataset, we partition it at the start with `partitionBy`, so that it does not need to be shuffled across the network. In practice, the `links` RDD is also likely to be much larger in terms of bytes than `ranks`, since it contains a list of neighbors for each page ID instead of just a `Double`, so this optimization saves considerable network traffic over a simple implementation of PageRank (e.g. in plain MapReduce).
2. For the same reason, we call `persist` on `links` to keep it in RAM across iterations.
3. When we first create `ranks`, we use `mapValues` instead of `map` to preserve the partitioning of the parent RDD (`links`), so that our first join against it is very cheap.
4. In the loop body, we follow our `reduceByKey` with `mapValues`; because the result of `reduceByKey` is already hash-partitioned, this will make it more efficient to join the mapped result against `links` on the next iteration.

Finally, note also that the extra syntax from using partitioning is small, and `mapValues` in particular is more concise in the places it's used here than a `map`.

## Note

To maximize the potential for partitioning-related optimizations, you should use `mapValues` or `flatMapValues` whenever you are not changing an element's key.

## Custom Partitioners

While Spark's `HashPartitioner` and `RangePartitioner` are well-suited to many use cases, Spark also allows you to tune how an RDD is partitioned by providing a custom `Partitioner` object. This can be used to further reduce communication by taking advantage of domain-specific knowledge.

For example, suppose we wanted to run the PageRank algorithm in the previous section on a set of web pages. Here each page's ID (the key in our RDD) will be its URL. Using a simple hash function to do the partitioning, pages with similar URLs (e.g., `http://www.cnn.com/WORLD` and `http://www.cnn.com/US`) might be hashed to completely different nodes. However, we know that web pages within the same domain tend to link to each other a lot. Because PageRank needs to send a message from each page to each of its neighbors on each iteration, it helps to group these pages into the same partition. We can do this with a custom `Partitioner` that looks at just the domain name instead of the whole URL.

To implement a custom partitioner, you need to subclass the `spark.Partitioner` class and implement three methods:

- `numPartitions: Int`, which returns the number of partitions you will create
- `getPartition(key: Any): Int`, which returns the partition ID (0 to `numPartitions-1`) for a given key
- `equals`, the standard Java equality method. This is important to implement because Spark will need to test your `Partitioner` object against other instances of itself when it decides whether two of your RDDs are partitioned in the same way!

One gotcha is that, if you rely on Java's `hashCode` method in your algorithm, it can return negative numbers. You need to be careful that `getPartition` always returns a non-negative result.

For example, here is how we would write the domain-name based partitioner sketched above, which hashes only the domain name of each URL:

```
class DomainNamePartitioner(numParts: Int) extends Partitioner {  
  override def numPartitions: Int = numParts  
  
  override def getPartition(key: Any): Int = {  
    val domain = new java.net.URL(key.toString).getHost()  
    val code = (domain.hashCode % numPartitions)  
    if (code < 0) {  
      code + numPartitions // Make it non-negative  
    } else {  

```

```

        code
    }
}

// Java equals method to let Spark compare our Partitioner objects
override def equals(other: Any): Boolean = other match {
    case dnp: DomainNamePartitioner =>
        dnp.numPartitions == numPartitions
    case _ =>
        false
}
}

```

Note that in the equals method, we used Scala's pattern matching operator (match) to test whether other is a DomainNamePartitioner, and cast it if so; this is the same as using instanceof in Java.

Using a custom Partitioner is easy: just pass it to the partitionBy method. Many of the shuffle-based methods in Spark, such as join and groupByKey, can also take an optional Partitioner object to control the partitioning of the output.

Creating a custom Partitioner in Java is very similar to Scala: just extend the spark.Partitioner class and implement the required methods.

In Python, you do not extend a Partitioner class, but instead pass a hash function as an additional argument to RDD.partitionBy(). For example:

```

import urlparse

def hash_domain(url):
    return hash(urlparse.urlparse(url).netloc)

rdd.partitionBy(20, hash_domain)    # Create 20 partitions

```

Note that the hash function you pass will be compared *by identity* to that of other RDDs. If you want to partition multiple RDDs with the same partitioner, pass the same function object (e.g., a global function) instead of creating a new lambda for each one!

## Conclusion

In this chapter, we have seen how to work with keyed data using the specialized functions available in Spark. The techniques from the [previous chapter on Programming with RDDs](#) also still work on our Pair RDDs. In the next chapter, we will look at how to load and save data.

---

[9] The Python API does not yet offer a way to query partitioners, though it still uses them internally.

## Chapter 5. Loading and Saving Your Data

Both engineers and data scientists will find parts of this chapter useful. Engineers may wish to explore more output formats to see if there is something well suited to their intended downstream consumer available and should consider looking online for different Hadoop formats. Data Scientists can likely focus on the input format that their data is already in. Spark supports reading from classes that implement Hadoop's `InputFormat` and writing to Hadoop's `OutputFormat` interfaces.

### Motivation

We've looked at a number of operations we can perform on our data once we have it distributed in Spark. So far our examples have loaded and saved all of their data from a native collection and regular files, but odds are that your data doesn't fit on a single machine, so it's time to explore our options.

In addition to different formats, we also have different compression options which can decrease the amount of space and network overhead required but can introduce restrictions on how we read the data. There are myriad of different data sources and formats we can use to create a new RDD and we will only cover the most common here.

### Note

In Spark 1.0, the standard data APIs for Python only support text data. However, you can work around this by using Spark SQL, which allows loading data from a number of formats supported by Hive. We will cover Spark SQL in a later chapter, but briefly illustrate how to load data from Hive in this chapter too.

### Choosing a Format

Often when we need data for our code we don't get to choose the format. Sometimes we are lucky enough to be able to choose from multiple formats or work with our upstream data providers. We will start with some [common file formats](#) which we can write to a number of different file systems. In addition to standard file formats, we will also examine using [different database systems](#) for IO. Spark works with all the formats implementing Hadoop's `InputFormat` and `OutputFormat` interfaces.

Table 5-1. Common Supported File Formats

Format name	Splittable	Structured	comments
text files	yes	no	Plain old text files. Splittable provided records are one per line.
JSON	yes	semi	Common text based format, are semi-structured, splittable if one record per line.
CSV	yes	yes	Very common text based format, often used with spreadsheet applications.
Sequence files	yes	yes	A common Hadoop file format used for key-value data.
Protocol Buffers	yes	yes	A fast space-efficient multi-language format.
Object Files	yes	yes	Useful for saving data from a Spark job to be consumed by shared code. Breaks if you change your classes, as it relies on Java Serialization.

In addition to these file formats, Spark can also work directly with a number of database and search systems. We will look at Mongo, Hive, Cassandra, and Elasticsearch, but many more are supported through the standard Hadoop connectors with Spark. This is convenient if your data happens to already be in one of these systems as nightly dumps of data can be difficult to maintain. When about to run a new Spark job against an on-line system you should verify you have sufficient capacity for a potentially very high volume of queries.

## Formats

Spark makes it very simple to load and save data in a large number of formats. Formats range from unstructured, like text, to semi-structured, like JSON, and to structured like Sequence Files. The input formats that Spark wraps all transparently handle compressed formats based on the file extension.

In addition to the output mechanism supported directly in Spark, we can use both Hadoop's new and old file APIs for keyed (or paired) data. This restriction exists because the Hadoop interfaces require key-value data, although some formats ignore the key.



## Text Files

Text files are very simple to load from and save to with Spark. When we load a single text file as an RDD each input line becomes an element in the RDD. We can also load multiple whole text files at the same time into a Pair RDD with the key being the name and the value being the contents of each file.

Loading a single text file is as simple as calling the `textFile` function on our spark context with the path to the file. If we want to control the number of partitions that the we can also specify `minPartitions`.

Example 5-1. Python load text file example

```
input = sc.textFile("file:///home/holden/repos/spark/README.md")
```

Example 5-2. Scala load text file example

```
val input = sc.textFile("file:///home/holden/repos/spark/README.md")
```

Example 5-3. Java load text file example

```
JavaRDD<String> input = sc.textFile("file:///home/holden/repos/spark/README.md")
```

Multi-part inputs in the form of a directory containing all of the parts can be handled in two ways. We can just use same `textFile` method and pass it a directory and it will load all of the parts into our RDD. Sometimes its important to know which file which piece of input came from (such as time data with the key in the file) or we need to process an entire file at a time. If our files are small enough then we can use the `wholeTextFiles` method and get back a Pair RDD where the key is the name of the input file.

`wholeTextFiles` can be very useful when each file represents a certain time periods data. If we had files representing sales data from different periods we could easily compute the average for each period.

Example 5-4. Scala average value per file

```
val input = sc.wholeTextFiles("file:///home/holden/happy panda")
val result = input.mapValues{y =>
  val nums = y.split(" ").map(x => x.toDouble)
  nums.sum / nums.size.toDouble
}
```

## Tip

Spark supports reading all the files in a given directory and doing wild card expansion on the input (e.g. `part-*.txt`). This is useful since large data sets are often spread across multiple files.

Outputting text files is also quite simple. The method `saveAsTextFile` takes a path and will output the contents of the RDD to that file. The path is treated as a directory and Spark will output multiple files underneath that directory. This allows Spark to write the output from multiple nodes. With this method we don't get to control which files end up with which segments of our data but there are other output formats which do allow this.

Example 5-5. Scala save as text file example

```
result.saveAsTextFile(outputFile)
```

Loading and saving text files is implemented through wrappers around Hadoop's file APIs. If you want further examples of how to work with the Hadoop file APIs you can see how this is implemented in `SparkContext` and `RDD` respectively.

## JSON

JSON is a popular semi-structured data format. The simplest way to load JSON data is by loading the data as a text file and then mapping over the values with a JSON parser. Like wise, we can use our preferred JSON serialization library to write out the values to strings which we can then write out. In Java and Scala we can also work with JSON data using a [custom Hadoop format](#).

Loading the data as a text file and then parsing the JSON data is an approach that we can use in all of the supported languages. This works assuming that you have one JSON record per-row, if you have multi-line JSON files you will instead have to load the whole file and then parse each file. If constructing a JSON parser is expensive in your language, you can use `mapPartitions` to re-use the parser.

There are a wide variety of JSON libraries available for the three languages we are looking at, for simplicities sake we are only considering one library per language. In Java we will use [Jackson](#), in Scala we are using [liftweb-json](#), and in python we use [the built in library](#). These libraries have been chosen as they perform reasonable well and are also relatively simple, if you spend a lot of time in the parsing stage you may wish to look at other JSON libraries [for Scala](#) or [for Java](#).

Example 5-6. Python load JSON example

```
data = input.map(lambda x: json.loads(x))
```

Example 5-7. Scala load JSON example

```
import play.api.libs.json._
import play.api.libs.functional.syntax._
...
val parsed = input.map(Json.parse(_))
```

If your JSON data happens to follow a predictable schema (lucky you!), we can parse it into a more structured format. This is often where we will find invalid records, so we may wish to skip them.

Example 5-8. Scala load JSON example

```
case class Person(name: String, lovesPandas: Boolean)
implicit val personReads = Json.format[Person]
// We use asOpt combined with flatMap so that if it fails to parse we
// get back a None and the flatMap essentially skips the result.
val parsed = input.flatMap(record =>
  personReads.read(JSON.parseFull(record)).asOpt)
```

Example 5-9. Java load JSON example

```
public static class ParseJson implements FlatMapFunction<Iterator<String>, Person> {
    public Iterable<Person> call(Iterator<String> lines) throws Exception {
        ArrayList<Person> people = new ArrayList<Person>();
        ObjectMapper mapper = new ObjectMapper();
        while (lines.hasNext()) {
            String line = lines.next();
            try {
                people.add(mapper.readValue(line, Person.class));
            } catch (Exception e) {
                // skip records on failure
            }
        }
        return people;
    }
}
```

## Tip

Handling incorrectly formatted records can be a big problem, especially with semi-structured data like JSON. With small data sets it can be acceptable to stop the world (i.e. fail the program) on malformed input, but often with large data sets malformed input is simply a part of life. If you do choose to skip incorrectly formatted (or attempt to recover) incorrectly formatted data you may wish to look at using accumulators to keep track of the number of errors.

Writing out JSON files is much simpler compared to loading it as we don't have to worry about incorrectly formatted data and we know the type of the data that we are writing out. We can use the same libraries we used to convert our RDD of strings into parsed JSON data and instead

take our RDD of structured data and convert it into an RDD of strings which we can then write out using Spark's text file API.

Lets say we were running a promotion for people that love pandas, so we can take our input from the first step and filter it for the people that love pandas.

Example 5-10. Python save JSON example

```
(data.filter(lambda x: x['lovesPandas'])).map(lambda x: json.dumps(x))
    .saveAsTextFile(outputFile))
```

Example 5-11. Scala save JSON example

```
result.filter(x => x.lovesPandas).map(x => Json.toJson(x)).
    saveAsTextFile(outputFile)
```

Example 5-12. Java save JSON example

```
public static class WriteJson implements
    FlatMapFunction<Iterator<Person>, String> {
    public Iterable<String> call(Iterator<Person> people) throws Exception {
        ArrayList<String> text = new ArrayList<String>();
        ObjectMapper mapper = new ObjectMapper();
        while (people.hasNext()) {
            Person person = people.next();
            text.add(mapper.writeValueAsString(person));
        }
        return text;
    }
}
```

```
JavaRDD<Person> result = input.mapPartitions(new ParseJson()).filter(
    new LikesPandas());
JavaRDD<String> formatted = result.mapPartitions(new WriteJson());
formatted.saveAsTextFile(outfile);
```

We can easily load and save JSON data with Spark by using the existing mechanism for working with text and adding JSON libraries.

## **CSV (Comma Separated Values) / TSV (Tab Separated Values)**

CSV files are supposed to contain a fixed number of fields per-line, and the fields are most commonly separated by comma or tab. Records are often stored one per line, but this is not always the case as records can sometimes span lines. CSV/TSV files can sometimes be inconsistent, most frequently in respect to handling newlines, escaping, non-ASCII characters,

non-integer numbers. CSVs cannot handle nested field types natively, so we have to unpack and pack to specific fields manually.

Unlike with JSON fields each record doesn't have field names associated with them; instead we get back row numbers. It is common practice in single CSV files to have the first rows column values be the names of each field.

Loading CSV/TSV data is similar to loading JSON data in that we can first load it as text and then process it. The lack of standardization leads to different versions of the same library sometimes handling input in different ways.

Like with JSON there are many different CSV libraries and we will only use one for each language. In both Scala and Java we use [opencsv](#). Once again in Python we use the included [csv](#) library.

## Tip

As with JSON there is a Hadoop [CSVInputFormat](#) that we can use to load CSV data in Scala and Java (although it does not support records containing newlines).

If your CSV data happens to not contain newlines in any of the fields, you can load your data with `textFile` and parse it.

Example 5-13. Python load CSV example

```
import csv
import StringIO
...
def loadRecord(line):
    """Parse a CSV line"""
    input = StringIO.StringIO(line)
    reader = csv.DictReader(input, fieldnames=["name", "favouriteAnimal"])
    return reader.next()
input = sc.textFile(inputFile).map(loadRecord)
```

Example 5-14. Java load CSV example

```
import au.com.bytecode.opencsv.CSVReader;
...
public static class ParseLine implements Function<String, String[]> {
    public String[] call(String line) throws Exception {
        CSVReader reader = new CSVReader(new StringReader(line));
        return reader.readNext();
    }
}
```

```

    }
}
JavaRDD<String> csvFile1 = sc.textFile(csv1);
JavaPairRDD<String[]> csvData = csvFile1.map(new ParseLine());

```

Example 5-15. Scala load CSV example

```

import au.com.bytecode.opencsv.CSVReader
...
val input = sc.textFile(inputFile)
val result = input.map{ line =>
    val reader = new CSVReader(new StringReader(line));
    reader.readNext();
}

```

If there are embedded newlines in fields we will need to load each file in full and parse the entire segment. This is unfortunate as if each file is large this can easily introduce bottlenecks in loading and parsing.

Example 5-16. Python load CSV example

```

def loadRecords(fileNameContents):
    """Load all the records in a given file"""
    input = StringIO.StringIO(fileNameContents[1])
    reader = csv.DictReader(input, fieldnames=["name", "favoriteAnimal"])
    return reader
fullFileData = sc.wholeTextFiles(inputFile).flatMap(loadRecords)

```

Example 5-17. Java load CSV example

```

public static class ParseLine implements FlatMapFunction<Tuple2<String, String>, String[]> {
    public Iterable<String[]> call(Tuple2<String, String> file) throws Exception {
        CSVReader reader = new CSVReader(new StringReader(file._2()));
        return reader.readAll();
    }
}

```

```

JavaPairRDD<String, String> csvData = sc.wholeTextFiles(csvInput);
JavaRDD<String[]> keyedRDD = csvData.flatMap(new ParseLine());

```

Example 5-18. Scala load CSV example

```

case class Person(name: String, favoriteAnimal: String)

val input = sc.wholeTextFiles(inputFile)
val result = input.flatMap{ case (_, txt) =>
    val reader = new CSVReader(new StringReader(txt));

```

```

    reader.readAll().map(x => Person(x(0), x(1)))
}

```

## Tip

If there are only a few input files you may want to repartition your input to allow Spark to effectively parallelize your future operations.

As with JSON data, writing out CSV/TSV data is quite simple and we can benefit from reusing the output encoding object. Since in CSV we don't output the field name with each record, to have a consistent output we need to create a mapping. One of the easy ways to do this is to just write a function which converts the fields to given positions in an array. In Python if we are outputting dictionaries the csv writer can do this for us based on the order we provide the fieldnames when constructing the writer.

The CSV libraries we are using output to files/writers so we can use `StringWriter/StringIO` to allow us to put the result in our RDD.

Example 5-19. Python write csv example

```

def writeRecords(records):
    """Write out CSV lines"""
    output = StringIO.StringIO()
    writer = csv.DictWriter(output, fieldnames=["name", "favoriteAnimal"])
    for record in records:
        writer.writerow(record)
    return [output.getvalue()]

```

```

pandaLovers.mapPartitions(writeRecords).saveAsTextFile(outputFile)

```

Example 5-20. Scala write CSV example

```

pandaLovers.map(person => List(person.name, person.favoriteAnimal)).toArray()
.mapPartitions{people =>
    val stringWriter = new StringWriter();
    val csvWriter = new CSVWriter(stringWriter);
    csvWriter.writeAll(people.toList)
    Iterator(stringWriter.toString)
}.saveAsTextFile(outFile)

```

As you may have noticed the above only works provided that we know all of the fields that we will be outputting. However, if some of the field names are determined at runtime from user

input we need to take a different approach. The simplest approach is going over all of our data and extracting the distinct keys.

## Sequence Files

Sequence files are a popular Hadoop format comprised of flat files with key-value pairs and are supported in Spark's Java and Scala APIs. Sequence files have sync markers that allow Spark to seek to a point in the file and then resynchronize with the record boundaries. This allows Spark to efficiently read Sequence files in from multiple nodes and in to many partitions. Sequence Files are a common input/output format for Hadoop MapReduce jobs as well so if you are working with an existing Hadoop system there is a good chance your data will be available as a sequence file.

Sequence files consist of elements which implement Hadoop's Writable interface, as Hadoop uses a custom serialization framework. We have a [conversion table](#) of some common types and their corresponding Writable class. The standard rule of thumb is try adding the word Writable to the end of your class name and see if it is a known subclass of [org.apache.hadoop.io.Writable](#). If you can't find a Writable for the data you are trying to write out (like for example a custom case class), you can go ahead and implement your own Writable class by overriding readFields and write from [org.apache.hadoop.io.Writable](#).

## Warning

Hadoop's RecordReader re-uses the same object for each record, so directly calling cache, on an RDD you read in like this can fail, instead add a simple map operation and cache the result of the map. Further more, many Hadoop Writable classes do not implement java.io.Serializable so for them to work in RDDs we need to convert them with a map anyways.

Table 5-2. Corresponding Hadoop Writable Types

Scala Type	Java Type	Hadoop Writable
Int	Integer	IntWritable or VIntWritable <a href="#">[a]</a>
Long	Long	LongWritable or VLongWritable <a href="#">[a]</a>
Float	Float	FloatWritable
Double	Double	DoubleWritable
Boolean	Boolean	BooleanWritable
Array[Byte]	Byte[]	BytesWritable
String	String	Text
Array[T]	T[]	ArrayWritable<TW> <a href="#">[b]</a>
List[T]	List<T>	ArrayWritable<TW> <a href="#">[b]</a>



Scala Type	Java Type	Hadoop Writable
Map[A, B]	Map<A, B>	MapWritable<AW, BW> <a href="#">[b]</a>
<a href="#">[a]</a> ints and longs are often stored as a fixed size. Storing the number 12 takes the same amount of space as storing the number 2**30. If you might have a large number of small numbers. Instead we can use variable sized types which will use less bits to store smaller numbers.		
<a href="#">[b]</a> The templated type must also be a writable type.		

Spark has a specialized API for reading in sequence files. On the Spark Context we can call `sequenceFile(path, keyClass, valueClass, minPartitions)`. As mentioned earlier, Sequence Files work with Writable classes, so our `keyClass` and `valueClass` will both have to be the correct Writable class. Lets consider loading people and the number of pandas they have seen from a sequence file, in this case our `keyClass` would be `Text` and our `valueClass` would be `IntWritable` or `VIntWritable`, for simplicity lets work with `IntWritable`.

Example 5-21. Scala load sequence file example

```
val data = sc.sequenceFile(inFile, classOf[Text], classOf[IntWritable]).
  map{case (x, y) => (x.toString, y.get())}
```

Example 5-22. Java load sequence file example

```
public static class ConvertToNativeTypes implements
  PairFunction<Tuple2<Text, IntWritable>, String, Integer> {
  public Tuple2<String, Integer> call(Tuple2<Text, IntWritable> record) {
    return new Tuple2(record._1.toString(), record._2.get());
  }
}
```

```
JavaPairRDD<Text, IntWritable> input = sc.sequenceFile(fileName, Text.class,
  IntWritable.class);
JavaPairRDD<String, Integer> result = input.mapToPair(
  new ConvertToNativeTypes());
```

## Tip

In Scala there is a convenience function which can automatically convert Writables to their corresponding Scala type. Instead of specifying the `keyClass` and `valueClass` we can call `sequenceFile[Key, Class](path, minPartitions)` and get back an RDD of native Scala types.

Writing the data out to a sequence file is fairly similar in Scala. First since sequence files are key-value pairs, we need a `PairRDD` with types that our sequence file can write out. Implicit conversions between Scala types and Hadoop Writables exist for many native types, so if you are

writing out a native type you can just save your PairRDD by calling `saveAsSequenceFile(path)` and it will write out the data for us. If there isn't an automatic conversion from our key and value to Writable, or we want to use VarInts we can just map over the data and convert it before saving. Lets consider writing out the data that we loaded in the previous example (people and how many pandas they have seen).

Example 5-23. Scala save sequence file example

```
val data = sc.parallelize(List(("Panda", 3), ("Kay", 6), ("Snail", 2)))
data.saveAsSequenceFile(outputFile)
```

In Java saving a sequence file is slightly more involved, due to the lack of `saveAsSequenceFile` method on the `JavaPairRDD`. Instead we use Spark's ability to save to [custom Hadoop formats](#) and we will show how to save to a sequence file in java in the custom Hadoop formats subsection.

## Object Files

Object files are a deceptively simple wrapper around sequence files which allows us to save our RDDs containing just values. Unlike with Sequence files, the values are written out using Java Serialization.

### Warning

If you change your classes, e.g., to add and remove fields, old object files may no longer be readable. Object files use Java Serialization, which has some support for managing compatibility across class versions but requires programmer effort to do so.

Using Java Serialization for object files has a number of implications. Unlike with normal sequence files, the output will be different than Hadoop outputting the same objects. Unlike the other formats, object files are mostly intended to be used for Spark jobs communicating with other Spark Jobs. Java Serialization can also be quite slow.

Saving an object file is as simple as calling `saveAsObjectFile` on an RDD. Reading an object file back is also quite simple, the function `objectFile` on the `SparkContext` takes in a path and returns an RDD.

With all of these warnings about object files you might wonder why anyone would use them. The primary reason to use object files are they require almost no work to save almost arbitrary objects.

## Hadoop Input and Output Formats

In addition to the formats Spark has wrappers for, we can also interact with other Hadoop supported formats. Spark supports both the old and new Hadoop file APIs providing a great amount of flexibility.

To read in a file using the new Hadoop API we need to tell spark a few things. The `newAPIHadoopFile` takes a path, and three classes. The first class is the “format” class, this is the class representing our input format. The next class is the class for our, and the final class is the class of our value. If we need to specify additional Hadoop configuration properties we can also pass in a conf object.

One of the simplest Hadoop input formats is the `KeyValueTextInputFormat` which can be used for reading in key-value data from text files. Each line is processed individually with they key and value separated by a tab character. This format ships with Hadoop so we don’t have to add any extra dependencies to our project to use it.

Example 5-24. Scala load `KeyValueTextInputFormat`

```
val input = sc.hadoopFile[Text, Text, KeyValueTextInputFormat](inputFile).map{
  case (x, y) => (x.toString, y.toString)
}
```

We looked at loading JSON data by loading the data as a text file and then parsing it, but we can also load JSON data using a custom Hadoop input format. This example requires setting up some extra bits for compression so feel free to skip it. Twitter’s [Elephant Bird package](#) supports a large number of data formats, including JSON, Lucene, Protocol Buffer related formats, and so on. The package also works with both the new and old Hadoop file APIs. To illustrate how to work with the new style Hadoop APIs from Spark lets look at loading LZO compressed JSON data with `LzoJsonInputFormat`:

Example 5-25. Scala load LZO compressed JSON with Elephant Bird

```
val input = sc.newAPIHadoopFile(inputFile, classOf[LzoJsonInputFormat],
  classOf[LongWritable], classOf[MapWritable], conf)
// Each MapWritable in "input" represents a JSON object
```

### Warning

LZO support requires installing the `hadoop-lzo` package and pointing Spark to its native libraries. If you install the Debian package, adding `--driver-library-path /usr/lib/hadoop/lib/native/ --driver-class-path /usr/lib/hadoop/lib/` to your spark-submit invocation should do the trick.

Reading a file using the old Hadoop API is pretty much the same from a usage point of view, except we provide an old style `InputFormat` class. Many of Spark's built in convenience functions (like `sequenceFile`) are implemented using the old style Hadoop API.

We already examined sequence files to some extent, but in Java we don't have the same convenience function for saving from a `PairRDD`. We will use this as a way to illustrate how to have using the old Hadoop format APIs as this is how Spark implements its helper function for `PairRDDs` in scala and we've already shown the new APIs with the JSON example.

Example 5-26. Java save sequence file example

public static class ConvertToWritableTypes implements

```
PairFunction<Tuple2<String, Integer>, Text, IntWritable> {
    public Tuple2<Text, IntWritable> call(Tuple2<String, Integer> record) {
        return new Tuple2(new Text(record._1), new IntWritable(record._2));
    }
}
```

```
JavaPairRDD<String, Integer> rdd = sc.parallelizePairs(input);
JavaPairRDD<Text, IntWritable> result = rdd.mapToPair(new ConvertToWritableTypes());
result.saveAsHadoopFile(fileName, Text.class, IntWritable.class,
    SequenceFileOutputFormat.class);
```

In addition to the `saveAsHadoopFile` and `saveAsNewAPIHadoopFile` functions, if you want more control over writing out a Hadoop format you can use `saveAsHadoopDataset` / `saveAsNewAPIHadoopDataset`. Both functions just take a configuration object on which you need to set all of the Hadoop properties. The configuration is done the same as one would do for configuring the output of a Hadoop MapReduce job.

## Protocol Buffers

[Protocol buffers](#) <sup>[10]</sup> were first developed at Google for internal RPCs and have since been open sourced. Protocol Buffers (PB) are structured data, with the fields and types of fields being clearly defined. Protocol Buffers are optimized to be fast for encoding and decoding and also take up the minimum amount of space. Compared to XML protocol buffers are 3x to 10x smaller and can be 20x to 100x faster to encode and decode. While a PB has a consistent encoding there are multiple ways to create a file consisting of many PB messages.

Protocol Buffers are defined using a domain specific language and then the protocol buffer compiler can be used to generate accessor methods in a variety of languages (including all those

supported by Spark). Since protocol buffers aim to take up a minimal amount of space they are not “self-describing” as encoding the description of the data would take up additional space. This means that to parse data which is formatted as PB we need the protocol buffer definition to make sense of data.

Protocol buffers consist of fields which can be either optional, required, or repeated. When parsing data, a missing optional field does not result in a failure, but a missing required field results in failing to parse the data. As such when adding new fields to existing protocol buffers it is good practice to make the new fields optional as not everyone will upgrade at the same time (and even if they do you might want to read your old data).

Protocol buffers fields can be many pre-defined types, or another protocol buffer message. These types include string, int32, enums, and more. This is by no means a complete introduction to protocol buffers, if you are interested you should consult the [protobuf website](#). For our example we will look at loading many VenueResponse objects from our sample proto.

Example 5-27. Sample protocol buffer definition

```
message Venue {
  required int32 id = 1;
  required string name = 2;
  required VenueType type = 3;
  optional string address = 4;

  enum VenueType {
    COFFEESHOP = 0;
    WORKPLACE = 1;
    CLUB = 2;
    OMNOMNOM = 3;
    OTHER = 4;
  }
}

message VenueResponse {
  repeated Venue results = 1;
}
```

Twitter’s Elephant Bird library, that we used in the previous section to load JSON data, also supports loading and saving data from protocol buffers. Lets look at writing out some Venues.

Example 5-28. Scala Elephant Bird Protocol buffer write out example

```
val job = new Job()
val conf = job.getConfiguration
LzoProtobufBlockOutputFormat.setClassConf(classOf[Places.Venue], conf);
val dnaLounge = Places.Venue.newBuilder()
dnaLounge.setId(1);
dnaLounge.setName("DNA Lounge")
dnaLounge.setType(Places.Venue.VenueType.CLUB)
val data = sc.parallelize(List(dnaLounge.build()))
val outputData = data.map{ pb =>
    val protoWritable = ProtobufWritable.newInstance(classOf[Places.Venue]);
    protoWritable.set(pb)
    (null, protoWritable)
}
outputData.saveAsNewAPIHadoopFile(outputFile, classOf[Text],
    classOf[ProtobufWritable[Places.Venue]],
    classOf[LzoProtobufBlockOutputFormat[ProtobufWritable[Places.Venue]]], conf)
```

## Tip

When building your project make sure to use the same protocol buffer library version as Spark, as of this writing that is version 2.5

## Hive and Parquet

A great way to create RDDs from Hive and Parquet is through using Spark SQL. This allows us to get back structured data from Hive and also write expressive queries. This approach also supports all three of our languages. We will cover this in detail the Spark SQL Chapter, but for now let's look at just loading data.

To connect Spark SQL to an existing Hive installation, you need to provide a Hive configuration. This is done by copying your hive-site.xml to Spark's conf/ directory. If you just want to explore, a local Hive metastore will be used if no hive-site.xml is set, and we can easily load data into a Hive table to query later on. Spark has an example file we can load into our a Hive table.

When loading data with Spark SQL, the resulting RDDs consist of Row objects. In Python you can treat the row object much like a hash map.

Example 5-29. Python Hive load example

```
from pyspark.sql import HiveContext
```

```
hiveCtx = HiveContext(sc)
rows = hiveCtx.hql("SELECT key, value FROM src")
keys = input.map(lambda row: row["key"])
```

In Java and Scala the Row objects we get back allow access based on the column number. Each Row object has a get method that gives back a general type we can cast, and specific get methods for common basic types (e.g. getFloat, getInt, getLong, getString, getShort, getBoolean).

Example 5-30. Scala Hive load example

```
import org.apache.spark.sql.hive.HiveContext

val hiveCtx = new org.apache.spark.sql.hive.HiveContext(sc)
val rows = hiveCtx.hql("SELECT key, value FROM src")
val keys = input.map(row => row.getInt(0))
```

Example 5-31. Java Hive load example

```
import org.apache.spark.sql.hive.api.java.JavaHiveContext;
import org.apache.spark.sql.api.java.Row;
import org.apache.spark.sql.api.java.JavaSchemaRDD;

JavaHiveContext hiveCtx = new JavaHiveContext(sc);
JavaSchemaRDD rows = hiveCtx.hql("SELECT key, value FROM src");
JavaRDD<Integer> keys = rdd.map(new Function<Row, Integer>() {
    public Integer call(Row row) { return row.getInt(0); }
});
```

When loading data from Hive, Spark SQL supports any Hive-supported storage format, including text files, RCFiles, ORC, Parquet, Avro and Protocol Buffers.

Without a Hive installation, Spark SQL can also directly load data from Parquet files:

Example 5-32. Python Parquet load example

```
from pyspark.sql import SQLContext

sqlCtx = SQLContext(sc)
rows = sqlCtx.parquetFile("people.parquet")
names = input.map(lambda row: row["name"])
```

## File Systems

Spark supports a large number of file systems for reading and writing to which we can use with any of the file formats we want.

### Local/"Regular" FS

While Spark supports loading files from the local file system, it can be somewhat less convenient to work with compared to the other options. While it doesn't require any setup, it **requires that the files are available on all the nodes in your cluster**.

Some network file systems, like NFS, AFS, MapR's NFS layer, are exposed to the user as a regular file system. If your data is already in one of these forms, then you can use it as an input by just specifying [pathToFile](#) as the path and Spark will handle it.

Example 5-33. Scala load compressed text file from local FS

```
val rdd = sc.textFile("file:///home/holden/happypandas.gz")
```

### Tip

If the file isn't already on all nodes in the cluster, you can load it locally and then call `parallelize`. We can also use `addFile(path)` to distribute the contents and then use `SparkFiles.get(path)` in place where we would normally specify the location (e.g. `sc.textFile(SparkFiles.get(...))`). Both of these approaches can be slow, so consider if you can put your files in HDFS, on S3 or similar.

### Amazon S3

S3 is an increasingly popular option for storing large amount of data. S3 is especially fast when our compute nodes are located inside of EC2, but can easily have much worse performance if we have to go over the public internet.

Spark will check the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` environment variables for your S3 credentials.

### HDFS

The Hadoop Distributed File System (HDFS) is a popular distributed file system that Spark works well with. HDFS is designed to work on commodity hardware and be resilient to node failure while providing high data throughput. Spark and HDFS can be collocated on the same machines and Spark can take advantage of this data locality to avoid network overhead.



Using Spark with HDFS is as simple as specifying `hdfs://master:port/path` for your input and output. The Deploying Spark chapter covers how to setup spark for HDFS systems requiring authentication.

## **Warning**

The HDFS protocol has changed, if you run a version of Spark which is compiled for a different version it will fail. By default Spark is build against 1.0.4, if you build from source you can specify `SPARK_HADOOP_VERSION=` as a environment variable to build against a different version or you can download a different precompiled version of Spark.

## **Compression**

Frequently when working with big data, we find our selves needing to use compressed data to save storage space and network overhead. With most Hadoop output formats we can specify a compression codec which will compress the data. As we have already seen, Spark's native input formats (`textFile` and `SequenceFile`) can automatically handle some types of compression for us. When reading in compressed data, there are some compression codecs which can be used to automatically guess the compression type.

These compression options only apply to the Hadoop formats which support compression, namely those which are written out to a file system. The database Hadoop formats generally do not implement support for compression, or if they have compressed records it is configured in the database its self rather than with the Hadoop connector.

Choosing an output compression codec can have a big impact on future users of the data. With distributed systems such as Spark we normally try and read our data in from multiple difference machines. To make this possible each worker needs to be able to find the start of a new record. Some compression formats make this impossible, which requires a single node read in all of the data which can easily lead to a bottleneck. Formats which can be easily read from multiple machines are called “splittable”.

Table 5-3. Compression Options

Format	Splittable	Average Compression Speed	Effectiveness on Text	Hadoop Compression Codec	Pure Java	Native	Comments
gzip	N	Fast	High	org.apache.hadoop.io.compress.GzipCodec	Y	Y	
lzo	Y <sup>[a]</sup>	Very fast	Medium	com.hadoop.compression.lzo.LzoCodec	Y	Y	LZO require installation on every worker node
bzip2	Y	Slow	Very high	org.apache.hadoop.io.compress.BZip2Codec	Y	Y	Uses pure Java for splittable version
zlib	N	Slow	Medium	org.apache.hadoop.io.compress.DefaultCodec	Y	Y	Default compression codec for Hadoop
Snappy	N	Very Fast	Low	org.apache.hadoop.io.compress.SnappyCodec	N	Y	There is a pure java port of Snappy but it is not currently available in Spark/Hadoop

<sup>[a]</sup> Depends on the library used

## Warning

While Spark's `textFile` method *can* handle compressed input, they automatically disable splittable even if the input is compressed in such a way that it could be read in a splittable way. If you find your self needing to read in a large single file compressed input, you should consider skipping Spark's wrapper and instead use either `newAPIHadoopFile` or `hadoopFile` and specify the correct compression codec.

Some data formats (like Sequence files) allow us to only compress the data of our key value data, which can be useful for doing lookups. Other data formats have their own compression control,

for example many of the formats in twitter's Elephant Bird package work with LZO compressed data.

Spark wraps both the old and new style APIs for specifying the compression codec.

If we don't know what the compression format is while writing our code, we can instead use the `CompressionCodecFactory` to determine the codec based on the file name.

## Databases

Spark load and write data with database, and database like systems in two primary ways. Spark SQL provides a query language and row interface for some databases. Additional databases, and database like systems, can be accessed through Hadoop's connectors.

## Elasticsearch

Spark can both read and write data from Elasticsearch using [ElasticSearch-Hadoop](#). Elasticsearch is a new open source Lucene based search system. Most of the connectors we have looked at so far have written out to files, this connector instead wraps RPCs to the Elasticsearch cluster.

The elastic search connector is a bit different than the other connectors we have examined, since it ignores the path information we provide instead depends on setting up configuration on our Spark context. The Elasticsearch `OutputFormat` connector also doesn't quite have the types to use Spark's wrappers, so we instead use `saveAsHadoopDataSet` which means we need to set more properties by hand. Lets look at how to read/write some simple data out to Elastic Search.

Example 5-34. Scala Elastic Search Output Example

```
val jobConf = new JobConf(sc.hadoopConfiguration)
jobConf.set("mapred.output.format.class", "org.elasticsearch.hadoop.mr.EsOutputFormat")
jobConf.setOutputCommitter(classOf[FileOutputCommitter])
jobConf.set(ConfigurationOptions.ES_RESOURCE_WRITE, "twitter/tweets")
jobConf.set(ConfigurationOptions.ES_NODES, "localhost")
FileOutputFormat.setOutputPath(jobConf, new Path("-"))
output.saveAsHadoopDataset(jobConf)
```

Example 5-35. Scala Elastic Search Input Example

```
def mapWritableToInput(in: MapWritable): Map[String, String] = {
  in.map{case (k, v) => (k.toString, v.toString)}.toMap
}
```

```

val jobConf = new JobConf(sc.hadoopConfiguration)
jobConf.set("mapred.output.format.class",
    "org.elasticsearch.hadoop.mr.EsOutputFormat")
jobConf.set(ConfigurationOptions.ES_RESOURCE_READ, args(1))
jobConf.set(ConfigurationOptions.ES_NODES, args(2))
val currentTweets = sc.hadoopRDD(jobConf,
    classOf[EsInputFormat[Object, MapWritable]], classOf[Object],
    classOf[MapWritable])
// Extract only the map
// Convert the MapWritable[Text, Text] to Map[String, String]
val tweets = currentTweets.map{ case (key, value) => mapWritableToInput(value) }

```

Compared to some of our other connectors this is a bit convoluted, but serves as a useful reference for how to work with these types of connectors.

## Warning

On the write side Elasticsearch can do mapping inference, but this can occasionally infer the types incorrectly, so it can be a good idea to [explicitly set a mapping](#) if you are storing things besides strings.

## Mongo

We will cover loading and saving data with Mongo and Spark in the next update of this book.

## Cassandra

Spark ships with examples of how to work with Cassandra and we will expand on this section in the future. For now if you want to use Spark with Cassandra, take a look at [the Cassandra examples in Spark](#).

## HBase

Spark ships with examples of how to work with HBase and we will expand on this section in the future. For now if you want to use Spark with HBase, take a look at [the HBase examples in Spark](#).

## Java Database Connectivity (JDBC)

In addition to using Hadoop input formats, you can create RDDs from JDBC queries. Unlike the other methods of loading data, rather than calling a method on the SparkContext we instead

create an instance of `org.apache.spark.rdd.JdbcRDD` and provide it with our `SparkContext` and the other input data it requires.

We will create a simple `JdbcRDD` using MySQL as

Example 5-36. Scala `JdbcRDD` Example

```
def createConnection() = {  
    Class.forName("com.mysql.jdbc.Driver").newInstance();  
    DriverManager.getConnection("jdbc:mysql://localhost/test?user=holden");  
}  
  
def extractValues(r: ResultSet) = {  
    (r.getInt(1), r.getString(2))  
}  
  
val data = new JdbcRDD(sc,  
    createConnection, "SELECT * FROM panda WHERE ? <= id AND ID <= ?",  
    lowerBound = 1, upperBound = 3, numPartitions = 2, mapRow = extractValues)  
println(data.collect().toList)
```

The min and max we provide to the `JdbcRDD` class allow Spark to query different ranges of the data on different machines, so we don't get bottlenecked trying to load all the data on a single node. If you don't know how many records there are, the you can just do a count query manually first and use the result.

Along similar lines we provide a function to establish the connection to our database. This lets each node create its own connection to load data over.

The last parameter converts the result from [java.sql.ResultSet](#) to a format that is useful for manipulating our data. If left out Spark will automatically convert everything to arrays of objects.

The `JdbcRDD` is currently accessed a bit differently from most of the other methods for loading data in to Spark, and provides another option for interfacing with database systems.

## Conclusion

With the end of this chapter you should be able to get your data into Spark to work with and store the result of your computation in a format that is useful for you. We have examined a number of different formats we can use for our data, as well as compression options and their

implications on how data can be consumed. Subsequent chapters will examine ways to write more effective and powerful Spark programs now that we can load and save large data sets.

---

[\[10\]](#) sometimes called pbs or protobufs

## About the Authors

Holden Karau is a software development engineer at Databricks and is active in open source. She is the author of an earlier Spark book. Prior to Databricks she worked on a variety of search and classification problems at Google, Foursquare, and Amazon. She graduated from the University of Waterloo with a Bachelors of Mathematics in Computer Science. Outside of software she enjoys paying with fire, welding, and hula hooping.

Most recently, Andy Konwinski co-founded Databricks. Before that he was a PhD student and then postdoc in the AMPLab at UC Berkeley, focused on large scale distributed computing and cluster scheduling. He co-created and is a committer on the Apache Mesos project. He also worked with systems engineers and researchers at Google on the design of Omega, their next generation cluster scheduling system. More recently, he developed and led the AMP Camp Big Data Bootcamps and first Spark Summit, and has been contributing to the Spark project.

Patrick Wendell is an engineer at Databricks as well as a Spark Committer and PMC member. In the Spark project, Patrick has acted as release manager for several Spark releases, including Spark 1.0. Patrick also maintains several subsystems of Spark's core engine. Before helping start Databricks, Patrick obtained an M.S. in Computer Science at UC Berkeley. His research focused on low latency scheduling for large scale analytics workloads. He holds a B.S.E in Computer Science from Princeton University

Matei Zaharia is a PhD student in the AMP Lab at UC Berkeley, working on topics in computer systems, cloud computing and big data. He is also a committer on Apache Hadoop and Apache Mesos. At Berkeley, he leads the development of the Spark cluster computing framework, and has also worked on projects including Mesos, the Hadoop Fair Scheduler, Hadoop's straggler detection algorithm, Shark, and multi-resource sharing. Matei got his undergraduate degree at the University of Waterloo in Canada.